

---

# Observational Robustness and Invariances in Reinforcement Learning via Lexicographic Objectives

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Policy robustness in Reinforcement Learning may not be desirable at any costs:  
2 the alterations caused by robustness requirements from otherwise optimal policies  
3 should be explainable, quantifiable and formally verifiable. In this work we study  
4 how policies can be *maximally robust* to arbitrary observational noise by analysing  
5 how they are altered by this noise through a stochastic linear operator interpretation  
6 of the disturbances, and establish connections between robustness and properties of  
7 the noise kernel and of the underlying MDPs. Then, we construct sufficient condi-  
8 tions for policy robustness, and propose a robustness-inducing scheme, applicable  
9 to any policy gradient algorithm, that formally trades off expected policy utility for  
10 robustness through *lexicographic optimisation*, while preserving convergence and  
11 sub-optimality of the original algorithm.

## 12 1 Introduction

13 Robustness in Reinforcement Learning (RL) [Morimoto and Doya, 2005] can be looked at from  
14 different perspectives: (1) distributional shifts in the training data with respect to the deployment  
15 stage Satia and Lave Jr [1973], Heger [1994], Nilim and El Ghaoui [2005], Xu and Mannor [2006];  
16 (2) uncertainty in the model or observations [Pinto et al., 2017, Everett et al., 2021]; (3) adversarial  
17 attacks against actions [Pattanaik et al., 2017, Fischer et al., 2019]; and (4) sensitivity of neural  
18 networks (used as policy or value function approximators) towards input disturbances [Kos and Song,  
19 2017, Huang et al., 2017]. Robustness does not naturally emerge in most RL settings, since agents  
20 are typically only trained in a single, unchanging environment: There is a trade-off between how  
21 robust a policy is and how close it is to the set of optimal policies in its training environment, and in  
22 safety-critical applications we may need to provide formal guarantees for this trade-off.

23 **Motivation** Consider a dynamical system where we need to synthesise a controller (policy) through  
24 a model-free approach. When using a simulator for training we expect the deployment of the controller  
25 in the real system to be affected by different sources of noise, possibly not predictable or modelled (*e.g.*  
26 for networked components we may have sensor faults, communication delays, *etc*). In safety-critical  
27 systems, robustness (in terms of successfully controlling the system under disturbances) should  
28 preserve formal guarantees, and plenty of effort has been put on developing formal convergence  
29 guarantees on policy gradient algorithms [Agarwal et al., 2021, Bhandari and Russo, 2019] which  
30 vanish when “robustifying” policies through regularisation or adversarial approaches. Therefore,  
31 for such applications one would need a scheme to regulate the robustness-utility trade-off in RL  
32 policies, that on the one hand preserves the formal guarantees of the original algorithms, and on the  
33 other attains sub-optimality conditions from the original problem. Additionally, if we do not know  
34 the structure of the disturbance (which holds in most applications), learning directly a policy for an  
35 arbitrarily disturbed environment will yield unexpected behaviours when deployed in the true system.

36 **Lexicographic Reinforcement Learning (LRL)** Recently, lexicographic optimisation [Isermann,  
 37 Rentmeesters et al., 1996] has been applied to the multi-objective RL setting [Skalse et al.,  
 38 2022b]. In an LRL setting with different reward-maximising objective functions  $\{K_i\}_{1 \leq i \leq n}$ , some  
 39 objectives may be more important than others, and so we may want to obtain policies that solve the  
 40 multi-objective problem in a lexicographically prioritised way, *i.e.*, “find the policies that optimise  
 41 objective  $i$  (reasonably well), and from those the ones that optimise objective  $i + 1$  (reasonably well),  
 42 and so on”. There exist both value- and policy-based algorithms for LRL, and the approach is broadly  
 43 applicable to (most) existing RL algorithms [Skalse et al., 2022b].

44 **Previous Work** In robustness against *model uncertainty*, the MDP may have noisy or uncertain  
 45 reward signals or transition probabilities, as well as possible resulting *distributional shifts* in the  
 46 training data [Heger, 1994, Xu and Mannor, 2006, Fu et al., 2018, Pattanaik et al., 2018, Pirotta  
 47 et al., 2013, Abdullah et al., 2019], which connects to ideas on distributionally robust optimisation  
 48 [Wiesemann et al., 2014, Van Parys et al., 2015]. One of the first examples is Heger [1994], where  
 49 the author proposes using minimax approaches to learn  $Q$  functions that minimise the worst case  
 50 total discounted cost in a general MDP setting. Derman et al. [2020] propose a Bayesian approach to  
 51 deal with uncertainty in the transitions. Another robustness sub-problem is studied in the form of  
 52 *adversarial attacks or disturbances* by considering adversarial attacks on policies or action selection  
 53 in RL agents [Gleave et al., 2020, Lin et al., 2017, Tessler et al., 2019, Pan et al., 2019, Tan et al.,  
 54 2020, Klima et al., 2019]. Recently, Gleave et al. [2020] propose the idea that instead of modifying  
 55 observations, one could attack RL agents by swapping the policy for an adversarial one at given  
 56 times. For a detailed review on Robust RL see Moos et al. [2022]. Our work focuses in the study of  
 57 robustness versus *observational disturbances*, where agents observe a disturbed state measurement  
 58 and use it as input for the policy [Kos and Song, 2017, Huang et al., 2017, Behzadan and Munir,  
 59 2017, Mandlekar et al., 2017, Zhang et al., 2020, 2021]. This problem emerges in many robotics  
 60 applications, where one learns a policy through a simulator or human imitation, and then needs  
 61 to rely on sensor data for a real-world deployment. In particular Mandlekar et al. [2017] consider  
 62 both random and adversarial state perturbations, and introduce physically plausible generation of  
 63 disturbances in the training of RL agents that make the resulting policy robust towards realistic  
 64 disturbances. Zhang et al. [2020] propose a *state-adversarial* MDP framework, and utilise adversarial  
 65 regularising terms that can be added to different deep RL algorithms to make the resulting policies  
 66 more robust to observational disturbances, and Zhang et al. [2021] study how LSTM increases  
 67 robustness with optimal state-perturbing adversaries.

## 68 1.1 Main Contributions

69 Most existing work on RL with observational disturbances proposes modifying RL algorithms (learning  
 70 to deal with perturbations through linear combinations of regularising loss terms or adversarial  
 71 terms) that come at the cost of *explainability* (in terms of sub-optimality bounds) and *verifiability*,  
 72 since the induced changes in the new policies result in a loss of convergence guarantees. Our main  
 73 contributions are summarised in the following points.

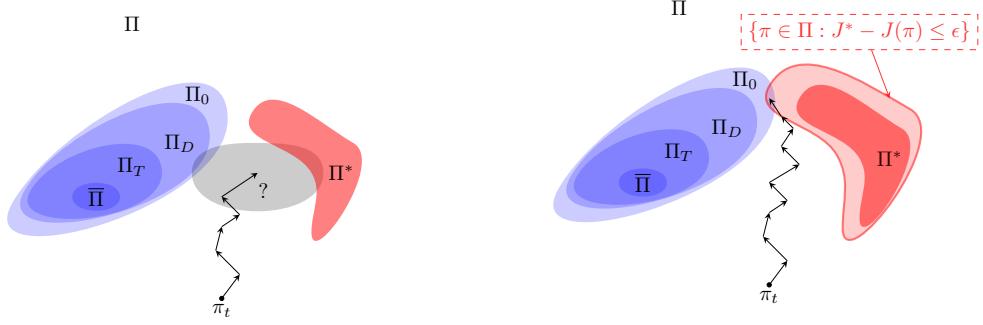
74 **Structure of Robust Policy Sets.** We consider general unknown stochastic disturbances and for-  
 75 mulate a quantitative definition of observational robustness that allows us to characterise the sets of  
 76 robust policies for any MDP in the form of operator-invariant sets. We analyse how the structure of  
 77 these sets depends on the MDP and noise kernel, and obtain an inclusion relation (cf. the Inclusion  
 78 Theorem, Section 3) providing intuition into how we can search for robust policies more effectively.<sup>1</sup>

79 **Verifiable Robustness through LRL.** The proposed characterisation and analysis allows us to  
 80 cast robustness as a lexicographic optimisation objective and propose a meta-algorithm that can be  
 81 applied to any existing policy gradient algorithm: Lexicographically Robust Policy Gradient (LRPG).  
 82 Compared to existing approaches for observational robustness, LRPG allows us to:

- 83     1. Retain policy sub-optimality up to a specified tolerance while maximising robustness.
- 84     2. Formally control the utility-robustness trade-off through this design tolerance.

---

<sup>1</sup>We claim novelty on the application of such concepts to the understanding and improvement of robustness in disturbed observation RL. Although we have not found our results in previous work, there are strong connections between Sections 2-3 in this paper and the literature on planning for POMDPs [Spaan and Vlassis, 2004, Spaan, 2012] and MDP invariances [Ng et al., 1999, van der Pol et al., 2020, Skalse et al., 2022a].



(a) PG algorithms when robustness terms are added to the cost function *indiscriminately*.

(b) In LRPG, the policy is guaranteed (up to the original algorithm used) to converge to an  $\epsilon$  ball of  $\Pi^*$ , and from those, the most robust ones.

Figure 1: Qualitative representation of the proposed LRPG algorithm, compared to usual robustness-inducing algorithms. The sets in blue are the maximally robust policies to be defined in the coming sections. Through LRPG we guarantee that the policies will only deviate a bounded distance from the original objective, and induce a search for robustness in the resulting valid policy set.

85        3. Preserve formal guarantees of the PG algorithm.

86        We provide numerical examples on how this approach is applied to existing policy gradient algorithms,  
 87        comparing them to previous work and verifying how the previously mentioned Inclusion Theorem  
 88        helps to induce more robust policies while retaining algorithm optimality. Figure 1 represents a  
 89        qualitative interpretation of the results in this work (the structure of the robust sets will become clear  
 90        in following sections).

91        1.2 Preliminaries

92        **Notation** We use calligraphic letters  $\mathcal{A}$  for collections of sets and  $\Delta(\mathcal{A})$  as the space of probability measures over  $\mathcal{A}$ . For two probability distributions  $P, P'$  defined on the same  $\sigma$ -algebra  $\mathcal{F}$ ,  $D_{TV}(P||P') = \sup_{A \in \mathcal{F}} |P(A) - P'(A)|$  is the total variation distance. For two elements of a vector space we use  $\langle \cdot, \cdot \rangle$  as the inner product. We use  $\mathbf{1}_n$  as a column-vector of size  $n$  that has all entries equal to 1. We say that an MDP is *ergodic* if for any policy the resulting Markov Chain (MC) is ergodic. We say that  $S$  is a  $n \times n$  row-stochastic matrix if  $S_{ij} \geq 0$  and each row of  $S$  sums to 1.

98        **Lexicographic Reinforcement Learning** We provide an introduction to Policy-Based Lexicographic RL (PB-LRL) for an example with two objective functions. Consider a parameterised 99 policy  $\pi_\theta$  with  $\theta \in \Theta$ , and two objective functions  $K_1$  and  $K_2$ . PB-LRL uses a multi-timescale 100 optimisation scheme to optimise  $\theta$  faster for higher-priority objectives, iteratively updating the 101 constraints induced by these priorities and encoding them via Lagrangian relaxation techniques 102 [Bertsekas, 1997]. Let  $\theta' \in \operatorname{argmax}_\theta K_1(\theta)$ . Then, PB-LRL can be used to find parameters 103  $\theta'' = \operatorname{argmax}_\theta K_2(\theta)$ , such that  $K_1(\theta) \geq K_1(\theta') - \epsilon$ . This is done through the estimated 104 gradient ascent update:  
 105

$$\theta \leftarrow \operatorname{proj}_\Theta [\theta + \nabla_\theta \hat{K}(\theta)], \quad \lambda \leftarrow \operatorname{proj}_{\mathbb{R}_{\geq 0}} [\lambda + \eta_t (\hat{k}_1 - \epsilon_t - K_1(\theta))], \quad (1)$$

106 where  $\hat{K}(\theta) := (\beta_t^1 + \lambda \beta_t^2) \cdot K_1(\theta) + \beta_t^2 \cdot K_2(\theta)$ ,  $\lambda$  is a Langrange multiplier,  $\beta_t^1, \beta_t^2, \eta_t$  are learning 107 rates<sup>2</sup>, and  $\hat{k}_1$  is an estimate of  $K_1(\theta')$ . Typically, we set  $\epsilon_t \rightarrow 0$ , though we can use other tolerances 108 too, e.g.,  $\epsilon_t = 0.9 \cdot \hat{k}_1$ . For more details on the convergence proofs and technicalities of PB-LRL we 109 refer the reader to Skalse et al. [2022b].

<sup>2</sup> We assume all learning rates in this work  $\alpha_t(x, u) \in [0, 1]$  ( $\beta_t, \eta_t$ ...) satisfy the conditions  $\sum_{t=1}^\infty \alpha_t(x, u) = \infty$  and  $\sum_{t=1}^\infty \alpha_t(x, u)^2 < \infty$ .

## 110 2 Observationally Robust Reinforcement Learning

111 Robustness-inducing methods in model-free RL must address the following dilemma: How do we deal  
 112 with uncertainty without an explicit mechanism to estimate such uncertainty during policy execution?  
 113 Consider an example of an MDP where, at policy roll-out phase, there is a non-zero probability of  
 114 measuring a “wrong” state. In such a scenario (even without adversarial uncertainty) optimal policies  
 115 can be almost useless: measuring the wrong state can lead to executing unboundedly bad actions.  
 116 This problem is represented by the following version of a noise-induced partially observable Markov  
 117 Decision Process [Spaan, 2012].

118 **Definition 2.1.** An observationally-disturbed MDP (DOMDP) is (a POMDP) defined by the tuple  
 119  $(X, U, P, R, T, \gamma)$  where  $X$  is a finite set of states,  $U$  is a set of actions,  $P : U \times X \mapsto \Delta(X)$   
 120 is a probability measure of the transitions between states and  $R : X \times U \times X \mapsto \mathbb{R}$  is a reward  
 121 function. The map  $T : X \mapsto \Delta(X)$  is a stochastic kernel induced by some unknown noise signal,  
 122 such that  $T(y | x)$  is the probability of measuring  $y$  while the true state is  $x$ , and acts only on the  
 123 state observations. At last  $\gamma \in [0, 1]$  is a reward discount.

124 In a DOMDP<sup>3</sup> agents can measure the full state, but the measurement will be disturbed by some  
 125 unknown random signal *in the policy deployment*. Unlike the POMDP setting, the agent has access to  
 126 the true state  $x$  during learning of the policies (i.e., the simulator is noise-free), but has no information  
 127 about the noise kernel  $T$  or a way to estimate it. The difficulty of acting in such DOMDP is that  
 128 the transitions are actually undisturbed and a function of the true state  $x$ , but agents will have to act  
 129 based on disturbed states  $\tilde{x} \sim T(\cdot | x)$ . We then need to construct policies that will be as robust as  
 130 possible against such noise, without being able to construct noise estimates. This setting, which is  
 131 distinguished from the POMDP one, reflects many robotic problems, where we can design a policy  
 132 for ideal noise-less conditions, and we know that at deployment there will likely be noise, data  
 133 corruption, adversarial perturbations, *etc.*, but we do not have a-priori knowledge on the structure  
 134 of this disturbance. A (memoryless) policy for the agent is a stochastic kernel  $\pi : X \mapsto \Delta(U)$ . For  
 135 simplicity, we overload notation on  $\pi$ , denoting by  $\pi(x, u)$  as the probability of taking action  $u$  at  
 136 state  $x$  under the stochastic policy  $\pi$  in the MDP, i.e.,  $\pi(x, u) = \Pr\{u | x\}$ . The value function of  
 137 a policy  $\pi$ ,  $V^\pi : X \mapsto \mathbb{R}$ , is given by  $V^\pi(x_0) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t R(x_t, \pi(x_t), x_{t+1})]$ . The action-value  
 138 function of  $\pi$  ( $Q$ -function) is given by  $Q^\pi(x, u) = \sum_{y \in X} P(x, u, y)(R(x, u, y) + \gamma V^\pi(y))$ . We  
 139 then define the objective function as  $J(\pi) := \mathbb{E}_{x_0 \sim \mu_0}[V^\pi(x_0)]$  with  $\mu_0$  being a distribution of initial  
 140 states, and we use  $J^* := \max_\pi J(\pi)$  and  $\pi^*$  as the optimal policy. If a policy is parameterised by  
 141  $\theta \in \Theta$  we write  $\pi_\theta$  and  $J(\theta)$ .

142 **Assumption 2.2.** For any DOMDP and policy  $\pi$ , the resulting MC is irreducible and aperiodic.

143 We now formalise a notion of *observational robustness*. Firstly, due to the presence of the stochastic  
 144 kernel  $T$ , the policy we are applying is altered as we are applying a collection of actions in a possibly  
 145 wrong state. This behaviour can be formally captured by:

$$\Pr\{u | x, \pi, T\} = \langle \pi, T \rangle(x, u) := \sum_{y \in X} T(y | x)\pi(y, u), \quad (2)$$

146 where  $\langle \pi, T \rangle : X \mapsto \Delta(U)$  is the *disturbed* policy, which averages the current policy given the error  
 147 induced by the presence of the stochastic kernel. Notice that  $\langle \cdot, T \rangle(x) : \Pi \mapsto \Delta(U)$  is an averaging  
 148 operator yielding the alteration of the policy due to noise. We can then define the *robustness regret*<sup>4</sup>:

$$\rho(\pi, T) := J(\pi) - J(\langle \pi, T \rangle). \quad (3)$$

149 **Definition 2.3 (Policy Robustness).** We say that a policy  $\pi$  is  $\kappa$ -*robust* against a stochastic kernel  $T$   
 150 if  $\rho(\pi, T) \leq \kappa$ . If  $\pi$  is 0-*robust* we say it is maximally robust. We define the sets of  $\kappa$ -robust policies,  
 151  $\Pi_\kappa := \{\pi \in \Pi : \rho(\pi, T) \leq \kappa\}$ , with  $\Pi_0$  being the set of maximally robust policies.

152 One can motivate the characterisation and models above from a control perspective, where policies  
 153 use as input discretised state measurements with possible sensor measurement errors. Formally  
 154 ensuring robustness properties when learning RL policies will, in general, force the resulting policies  
 155 to deviate from optimality in the undisturbed MDP. We propose then the following problem.

<sup>3</sup>Definition 2.1 is a generalised form of the State-Adversarial MDP used by Zhang et al. [2020]: the adversarial case is a particular form of DOMDP where  $T$  assigns probability 1 to one adversarial state.

<sup>4</sup>The robustness regret satisfies  $\rho(\pi^*, T) \geq 0 \forall T$ , and it allows us to directly compare the robustness regret with the utility regret of the policy.

156 **Problem 1.** For a DOMDP and a given tolerance level  $\epsilon$ , derive a policy  $\pi^\epsilon$  that satisfies  $J^* - J(\pi^\epsilon) \leq \epsilon$   
 157 as a prioritised objective and is as robust as possible according to Definition 2.3.

### 158 3 Characterisation of Robust Policies

159 An important question to be addressed, before trying to synthesise robust policies through LRL, is  
 160 what these robust policies look like, and how they are related to DOMDP properties. The robustness  
 161 notion in Definition 2.3 is intuitive and it allows us to classify policies. We begin by exploring what  
 162 are the types of policies that are maximally robust, starting with the set of constant policies and set of  
 163 fix point of the operator  $\langle \cdot, T \rangle$ , whose formal descriptions are now provided.

164 **Definition 3.1.** A policy  $\pi : X \mapsto \Delta(U)$  is said to be constant if  $\pi(x) = \pi(y)$  for all  $x, y \in X$ , and  
 165 the collection of all constant policies is denoted by  $\bar{\Pi}$ . A policy  $\pi : X \mapsto \Delta(U)$  is called a fixed  
 166 point of the operator  $\langle \cdot, T \rangle$  if  $\pi(x) = \langle \pi, T \rangle(x)$  for all  $x \in X$ . The collection of all fixed points will  
 167 be denoted by  $\Pi_T$ .

168 In other words, a constant policy is any policy that yields the same action distribution for any state,  
 169 and a fixed point policy is any policy whose action distributions are un-altered by the noise kernel.  
 170 Observe furthermore that  $\Pi_T$  only depends on the kernel  $T$  and the set<sup>5</sup>  $X$ . We now present a  
 171 proposition that links the two sets of policies in Definition 3.1 with our notion of robustness.

172 **Proposition 3.2.** Consider a DOMDP as in Definition 2.1, the robustness notion given in Definition  
 173 2.3 and the concepts in Definition 3.1, then we have that  $\bar{\Pi} \subseteq \Pi_T \subseteq \Pi_0$ .

174 The importance of Proposition 3.2 is that it allows us to produce (approximately) maximally robust  
 175 policies by computing the distance of a policy to either the set of constant policies or to the fix point  
 176 of the operator  $\langle \cdot, T \rangle$ , and this is at the core of the construction in Section 4. However, before this, let  
 177 us introduce another set that is sandwiched between  $\Pi_0$  and  $\Pi_T$ . Let us assume we have a policy  
 178 iteration algorithm that employs an action-value function  $Q^\pi$  and policy  $\pi$ . The advantage function  
 179 for  $\pi$  is defined as  $A^\pi(x, u) := Q^\pi(x, u) - V^\pi(x)$  and can be used as a maximisation objective  
 180 to learn optimal policies (as in, e.g., A2C [Sutton et al., 1999], A3C [Mnih et al., 2016]). We can  
 181 similarly define the *noise disadvantage* (a form of negative advantage) of policy  $\pi$  as:

$$D^\pi(x, T) := V^\pi(x) - \mathbb{E}_{u \sim \langle \pi, T \rangle(x)}[Q^\pi(x, u)], \quad (4)$$

182 which measures the difference of applying at state  $x$  an action according to the policy  $\pi$  with that  
 183 of playing an action according to  $\langle \pi, T \rangle$  and then continuing playing an action according to  $\pi$ . Our  
 184 intuition says that if it happens to be the case that  $D^\pi(x, T) = 0$  for all states in the DOMDP, then  
 185 such a policy is maximally robust. And this is indeed the case, as shown in the next proposition.

186 **Proposition 3.3.** Consider a DOMDP as in Definition 2.1 and the robustness notion as in Definition  
 187 2.3. If a policy  $\pi$  is such that  $D^\pi(x, T) = 0$  for all  $x \in X$ , then  $\pi$  is maximally robust, i.e., let

$$\Pi_D := \{\pi \in \Pi : \mu_\pi(x) D^\pi(x, T) = 0 \forall x \in X\},$$

188 then we have that  $\Pi_D \subseteq \Pi_0$ .

189 So far we have shown that both the set of fixed points  $\bar{\Pi}$  and the set of policies for which the  
 190 disadvantage function is equal to zero  $\Pi_D$  are contained in the set of maximally robust policies. More  
 191 interesting is the fact that the inclusion established in Proposition 3.2 and the one in Proposition 3.3  
 192 can be linked in a natural way through the following Inclusion Theorem.

193 **Theorem 3.4** (Inclusion Theorem). For a DOMDP with noise kernel  $T$ , consider the sets  $\bar{\Pi}, \Pi_T, \Pi_D$   
 194 and  $\Pi_0$ . Then, the following inclusion relation holds:

$$\bar{\Pi} \subseteq \Pi_T \subseteq \Pi_D \subseteq \Pi_0.$$

195 Additionally, the sets  $\bar{\Pi}, \Pi_T$  are convex for all MDPs and kernels  $T$ , but  $\Pi_D, \Pi_0$  may not be.

196 Let us reflect on the inclusion relations of Theorem 3.4. The inclusions are in general not strict, and in  
 197 fact the geometry of the sets (as well as whether some of the relations are in fact equalities) is highly

---

<sup>5</sup>There is a (natural) bijection between the set of constant policies and the space  $\Delta(U)$ . The set of fixed points of the operator  $\langle \cdot, T \rangle$  also has an algebraic characterisation in terms of the null space of the operator  $\text{Id}(\cdot) - \langle \cdot, T \rangle$ . We are not exploiting the later characterisation in this paper.

198 dependent on the reward function, and in particular on the complexity (from an information-theoretic  
 199 perspective) of the reward function. As an intuition, less complex reward functions (more uniform)  
 200 will make the inclusions above expand to the entire policy set, and more complex reward functions  
 201 will make the relations collapse to equalities. The following Corollary illustrates this.

202 **Corollary 3.5.** *For any ergodic DOMDP there exist reward functions  $\bar{R}$  and  $\underline{R}$  such that the resulting  
 203 DOMDP satisfies: (i)  $\Pi_D = \Pi_0 = \Pi$  (any policy is max. robust) if  $R = \bar{R}$ , (ii)  $\Pi_T = \Pi_D = \Pi_0$   
 204 (only fixed point policies are maximally robust) if  $R = \underline{R}$ .*

205 We can now summarise the insights from Theorem B.3 and Corollary 3.5 in the following conclusions:  
 206 (1) The set  $\bar{\Pi}$  is maximally robust, convex and *independent of the DOMDP*, (2) The set  $\Pi_T$  is  
 207 maximally robust, convex, includes  $\bar{\Pi}$ , and its properties *only depend on T*, (3) The set  $\Pi_D$  includes  
 208  $\Pi_T$  and is maximally robust, but its properties *depend on the DOMDP*.

## 209 4 Robustness through Lexicographic Objectives

210 We have now characterised robustness in a DOMDP and explored the relation between the sets of  
 211 policies that are robust according to the definition proposed. We have seen in the Inclusion Theorem  
 212 that several classes of policies are maximally robust, and our goal now is to connect these results with  
 213 lexicographic optimisation. To be able to apply LRL results to our robustness problem we need to first  
 214 cast robustness as a valid objective to be maximised, and then show that a stochastic gradient descent  
 215 approach would indeed find a global maximum of the objective, therefore yielding a maximally  
 216 robust policy. Then, this robustness objective can be combined with a primary reward-maximising  
 217 objective  $K_1(\theta) = \mathbb{E}_{x_0 \sim \mu_0}[V^{\pi_\theta}(x_0)]$  and any algorithm with certified convergence to solve Problem  
 218 1. Policy-based LRL (PB-LRL) allows us to encode the idea that, when learning how to solve an RL  
 219 task, robustness is important but *not at any price*, i.e., we would like to solve the original objective  
 220 reasonably well<sup>6</sup>, and from those policies efficiently find the most robust one.

### 221 4.1 Robustness Objectives

222 We propose now a valid lexicographic objective for which a minimising solution yields a maximally  
 223 robust policy. For this, we will perturb the policy during training according to the following logic. In  
 224 the introduction, we emphasised that the motivation for this work comes partially from the fact that  
 225 we may not know  $T$  in reality, or have a way to estimate it. However, the theoretical results until now  
 226 depend on  $T$ . Our proposed solution to this lies in the results of Theorem 3.4. We can use a *design*  
 227 generator  $\tilde{T}$  to perturb the policy during training such that  $\tilde{T}$  has the smallest possible fixed point set  
 228 (i.e. the constant policy set), and any algorithm that drives the policy towards the set of fixed points  
 229 of  $\tilde{T}$  will also drive the policy towards fixed points of  $T$ : from Theorem 3.4,  $\Pi_{\tilde{T}} \subseteq \Pi_T$ .

230 **Assumption 4.1.** The design kernel  $\tilde{T}$  satisfies  $\Pi_{\tilde{T}} = \bar{\Pi}$

231 We discuss further the choice and implications of using a design kernel  $\tilde{T}$  in Section 5. One of the  
 232 messages of the Inclusion Theorem is the fact that fixed point policies are maximally robust. Consider  
 233 the objective to be minimised:

$$K_{\tilde{T}}(\theta) = \sum_{x \in X} \mu_{\pi_\theta}(x) \frac{1}{2} \|\pi_\theta(x) - \langle \pi_\theta, \tilde{T} \rangle(x)\|_2^2, \quad (5)$$

234 Notice that optimising (5) projects the current policy onto the set of fixed points of the operator  $\langle \cdot, \tilde{T} \rangle$ ,  
 235 and due to Assumption 2.2, which requires  $\mu_{\pi_\theta}(x) > 0$  for all  $x \in X$ , the optimal solution is equal  
 236 to zero if and only if there exists a value of the parameter  $\theta$  for which the corresponding  $\pi_\theta$  is a  
 237 fixed point of  $\langle \cdot, \tilde{T} \rangle$ . In practice, the objectives are computed for a batch of trajectory sampled states  
 238  $X_s \subset X$ , and averaged over  $\frac{1}{|X_s|}$ ; we denote these approximations with a hat. By applying standard  
 239 stochastic approximation arguments, we can prove that convergence is guaranteed for a SGD iteration  
 240 using  $\nabla_\theta \hat{K}_{\tilde{T}}(\theta)(x) = (\pi_\theta(x) - \pi_\theta(y)) \nabla_\theta \pi_\theta(x)$ ,  $y \sim \tilde{T}(\cdot | x)$  to the optimal solution of problem 5.  
 241 For details and a proof, see Lemma B.3 in Appendix B.

---

<sup>6</sup>The advantage of using LRL is that we need not know in advance how to define “reasonably well” for each new task. Additionally, we obtain a hyper-parameter that directly controls the trade-off between *robustness and optimality*: the tolerance  $\epsilon$ . Through  $\epsilon$  we determine how far we allow our resulting policy to be from an optimal policy in favour of it being more robust.

242 **4.2 Lexicographically Robust Policy Gradient**

243 We present now the proposed LRPB meta-algorithm to achieve lexicographic robustness for any policy gradient algorithm at choice.  
 244 From Skalse et al. [2022b], the convergence of PB-LRL algorithms is guaranteed as long as  
 245 the original policy gradient algorithm (such as PPO [Liu et al., 2019] or A2C [Konda and Tsitsiklis, 2000, Bhatnagar et al., 2009]) for each single objective converges. We can then combine Lemma B.3 with these results to guarantee that Lexicographically Robust Policy Gradient (LRPG), Algorithm 1, converges to a policy that maximise robustness while remaining (approximately) optimal with respect to  $R$ .

256 **Theorem 4.2.** Consider a DOMDP as in Definition 2.1 and let  $\pi_\theta$  be a parameterised policy. Take  $K_1(\theta) = \mathbb{E}_{x_0 \sim \mu_0}[V^{\pi_\theta}(x_0)]$  to be computed through a chosen algorithm (e.g., A2C, PPO) that optimises  $K_1(\theta)$ , and let  $K_2(\theta) = -K_{\tilde{T}}(\theta)$ . Given an  $\epsilon > 0$ , if the iteration  $\theta \leftarrow \text{proj}_\Theta[\theta + \nabla_\theta \hat{K}_1]$  is guaranteed to converge to a parameter set  $\theta^*$  that maximises  $K_1$ , and hence  $J$  (locally or globally), then LRPG converges a.s. under PB-LRL conditions to parameters  $\theta^\epsilon$  that satisfy:

$$\theta^\epsilon \in \underset{\theta \in \Theta'}{\operatorname{argmin}} K_{\tilde{T}}(\theta), \quad \text{such that} \quad K_1^* \geq K_1(\theta^\epsilon) - \epsilon, \quad (6)$$

266 where  $\Theta' = \Theta$  if  $\theta^*$  is globally optimal and a compact local neighbourhood of  $\theta^*$  otherwise.

267 We reflect again on Figure 1. The main idea behind LRPG is that by formally expanding the set of acceptable policies with respect to  $K_1$ , we may find robust policies more effectively while guaranteeing a minimum performance in terms of expected rewards. This addresses directly the premise behind Problem 1. In LRPG the first objective is still to minimise the distance  $J^* - J(\pi)$  up to some tolerance. Then, from the policies that satisfy this constraint, we want to steer the learning algorithm towards a maximally robust policy, and we can do so without knowing  $T$ .

273 **5 Considerations on Noise Generators**

274 A natural question following Section 4.1 and the theoretical results in Section 4 is how to choose  $\tilde{T}$ , and how the choice influences the resulting policy robustness towards any other true  $T$ . In general, for any arbitrary policy utility landscape in a given MDP, there is no way of bounding the distance of the resulting policies for two different noise kernels  $T_1, T_2$ . As a counter-example, consider an MDP where there are 2 possible optimal policies  $\pi_1^*, \pi_2^*$ , and take these two policies to be maximally different, i.e.  $D_{TV}(\pi_1^* \| \pi_2^*) = 1 \forall x \in X$ . Then, when using LRPG to obtain a robust policy, a slight deviation in the choice of  $\tilde{T}$  can cause the gradient descent scheme to deviate from converging to  $\pi_1^*$  to converging to  $\pi_2^*$ , yielding in principle a completely different policy. However, the optimality of the policy remains bounded: Through LRPG guarantees we know that, for both cases, the utility of the resulting policy will be at most  $\epsilon$  far from the optimal. We can, thus, state the following.

284 **Corollary 5.1.** Take  $T$  to be any arbitrary noise kernel, and  $\tilde{T}$  to satisfy Assumption 4.1. Let  $\pi$  be a policy resulting from a LRPG algorithm. Assume that  $\min_{\pi' \in \Pi_{\tilde{T}}} D_{TV}(\pi \| \pi') = a$  for some  $a < 1$ . Then, it holds for any  $T$  that  $\min_{\pi' \in \Pi_T} D_{TV}(\pi \| \pi') \leq a$ .

287 That is, when using LRPG to obtain a robust policy  $\pi$ , the resulting policy is at most  $a$  far from the set of fixed points (and therefore a maximally robust policy) with respect to the true  $T$ . This is 288 the key argument behind our choices for  $\tilde{T}$ : A priori, the most sensible choice is a kernel that has 289 no other fixed point than the set of constant policies. This fixed point condition is satisfied in the 290 discrete state case for any  $\tilde{T}$  that induces an irreducible Markov Chain, and in continuous state for 291

PPO on MiniGrid Environments					A2C on MiniGrid Environments				
Noise	Vanilla	LR <sub>PPO</sub> ( $K_T^u$ )	LR <sub>PPO</sub> ( $K_T^g$ )	SA-PPO	Vanilla	LR <sub>A2C</sub> ( $K_T^u$ )	LR <sub>A2C</sub> ( $K_T^g$ )	LR <sub>A2C</sub> ( $K_D$ )	
<i>LavaGap</i>									
$\emptyset$	<b>0.95±0.003</b>	<b>0.95±0.075</b>	<b>0.95±0.101</b>	0.94±0.068	<b>0.94±0.004</b>	<b>0.94±0.005</b>	<b>0.94±0.003</b>	<b>0.94±0.006</b>	
$T_1$	0.80±0.041	<b>0.95±0.078</b>	0.93±0.124	0.88±0.064	0.83±0.061	<b>0.93±0.019</b>	0.89±0.032	0.91±0.088	
$T_2$	0.92±0.015	<b>0.95±0.052</b>	<b>0.95±0.094</b>	0.93±0.050	0.89±0.029	<b>0.94±0.008</b>	0.93±0.011	0.93±0.021	
<i>LavaCrossing</i>									
$\emptyset$	<b>0.95±0.023</b>	0.93±0.050	0.93±0.018	0.88±0.091	0.91±0.024	0.91±0.063	0.90±0.017	<b>0.92±0.034</b>	
$T_1$	0.50±0.110	<b>0.92±0.053</b>	0.89±0.029	0.64±0.109	0.66±0.071	<b>0.78±0.111</b>	0.72±0.073	0.76±0.098	
$T_2$	0.84±0.061	<b>0.92±0.050</b>	<b>0.92±0.021</b>	0.85±0.094	0.78±0.054	0.83±0.105	0.86±0.029	<b>0.87±0.063</b>	
<i>DynamicObstacles</i>									
$\emptyset$	<b>0.91±0.002</b>	<b>0.91±0.008</b>	<b>0.91±0.007</b>	<b>0.91±0.131</b>	<b>0.91±0.011</b>	0.88±0.020	0.89±0.009	<b>0.91±0.013</b>	
$T_1$	0.23±0.201	<b>0.77±0.102</b>	0.61±0.119	0.45±0.188	0.27±0.104	0.43±0.108	0.45±0.162	<b>0.56±0.270</b>	
$T_2$	0.50±0.117	<b>0.75±0.075</b>	0.70±0.072	0.68±0.490	0.45±0.086	0.53±0.109	0.52±0.161	<b>0.67±0.203</b>	

Table 1: Reward values gained by LRPG and baselines on discrete control tasks.

any  $\tilde{T}$  that satisfies a reachability condition (*i.e.* for any  $x_0 \in X$ , there exists a finite time for which the probability of reaching any ball  $B \subset X$  of radius  $r > 0$  through a sequence  $x_{t+1} = T(x_t)$  is measurable). This holds for (additive) uniform or Gaussian disturbances.

## 6 Experiments

We verify the theoretical results of LRPG in a series of experiments on discrete state/action safety-related environments [Chevalier-Boisvert et al., 2018], and in continuous control tasks. We use A2C [Sutton and Barto, 2018] (LR-A2C), PPO [Schulman et al., 2017] (LR-PPO) and SAC [Haarnoja et al., 2018] (LR-SAC) for our implementations of LRPG. In all cases, the lexicographic tolerance was set to  $\epsilon = 0.99\hat{k}_1$  to deviate as little as possible from the primary objective. We compare against the baseline algorithms and against SA-PPO [Zhang et al., 2021] which is among the most effective (adversarial) robust RL approaches in literature. We trained 10 independent agents for each algorithm, and reported the scores of the median agent (as done in Zhang et al. [2020]) for 50 averaged roll-outs.

**Sampling  $\tilde{T}$ .** To simulate  $\tilde{T}$  we disturb  $x$  as  $\tilde{x} = x + \xi$  for (1) a uniform bounded noise signal  $\xi \sim \mathcal{U}_{[-b,b]}(\tilde{T}^u)$  and (2) and a Gaussian noise ( $\tilde{T}^g$ ) such that  $\xi \sim \mathcal{N}(0, 0.5)$ . We test the resulting policies against a noiseless environment ( $\emptyset$ ), a kernel  $T_1 = \tilde{T}^u$ , a kernel  $T_2 = \tilde{T}^g$  and against two different state-adversarial noise configurations as proposed by Zhang et al. [2021] to evaluate how effective LRPG is at rejecting adversarial disturbances. See Appendix C for details and bounds used.

**Robustness Objectives.** If we do not have an estimator for the critic  $Q^\pi$  (*e.g.* PPO, A2C), Proposition 3.2 suggests that minimising the distance between  $\pi$  and  $\langle \pi, T \rangle$  can serve as a proxy to minimise the robustness regret, so we use objectives as defined in (5). We aim to test the hypothesis introduced through this work: If we have an estimator for the critic  $Q^\pi$  we can obtain robustness without inducing regularity in the policy using  $D^\pi$ , yielding a larger policy subspace to steer towards, and hopefully achieving policies closer to optimal. With the goal of diving deeper into the results of Theorem 3.4, we consider the objective  $K_D(\theta) := \sum_{x \in X} \mu_{\pi_\theta}(x) \frac{1}{2} \|D^{\pi_\theta}(x, T)\|_2^2$ . We use both in our experimental results, by modifying A2C to retain a Q critic.

**Robustness Results: Discrete Control.** Firstly, we investigate the impact of LRPG PPO and A2C for discrete action-space problems on Gymnasium [Brockman et al., 2016]. *Minigrid-LavaGap* (fully observable), *Minigrid-LavaCrossing* (partially observable) are safe exploration tasks where the agent needs to navigate an environment with cliff-like regions. *Minigrid-DynamicObstacles* (stochastic, partially observable) is a dynamic obstacle-avoidance environment. We use A2C to test the influence of  $K_D$  vs.  $K_T$  since the structure of the original cost functions are simpler than in PPO, and hence easier to compare between the scenarios above. With each objective function resulting in gradient descent steps that pull the policy towards different maximally robust sets ( $K_T \rightarrow \Pi_T$  and  $K_D \rightarrow \Pi_D$  respectively), we would expect to obtain increasing robustness for  $K_D$ . The results are presented in Table 1. See Appendix C for the results against adversarial noise, learning curves and detailed results.

PPO on Continuous Environments				SAC on Continuous Environments			
Noise	Vanilla	LR <sub>PPO</sub> ( $K_T^u$ )	LR <sub>PPO</sub> ( $K_T^g$ )	SA-PPO	Vanilla	LR <sub>SAC</sub> ( $K_T^u$ )	LR <sub>SAC</sub> ( $K_T^g$ )
<i>MountainCar</i>							
$\emptyset$	<b>94.77 ± 0.26</b>	93.17 ± 0.89	94.66 ± 1.61	88.69 ± 3.93	93.52 ± 0.05	<b>94.43 ± 0.19</b>	93.84 ± 0.05
$T_1$	88.67 ± 1.41	91.46 ± 1.22	<b>94.91 ± 1.35</b>	88.41 ± 3.99	1.89 ± 65.31	71.81 ± 13.04	<b>76.90 ± 7.11</b>
$T_2$	92.22 ± 1.11	92.40 ± 1.28	<b>94.76 ± 1.42</b>	89.32 ± 3.79	-27.82 ± 73.10	<b>72.93 ± 8.57</b>	69.41 ± 13.03
<i>LunarLander</i>							
$\emptyset$	267.99 ± 38.04	<b>269.76 ± 22.93</b>	243.08 ± 37.03	220.18 ± 98.78	268.96 ± 51.52	275.17 ± 14.04	<b>282.24 ± 15.95</b>
$T_1$	156.09 ± 22.87	<b>280.91 ± 20.34</b>	182.80 ± 49.26	164.53 ± 45.48	128.18 ± 17.73	<b>187.64 ± 76.30</b>	153.81 ± 33.16
$T_2$	158.02 ± 46.57	<b>276.76 ± 16.20</b>	212.62 ± 37.56	221.84 ± 73.61	140.92 ± 20.61	<b>187.82 ± 25.27</b>	158.18 ± 28.60
<i>BipedalWalker</i>							
$\emptyset$	265.39 ± 82.36	261.39 ± 83.19	<b>276.66 ± 44.85</b>	251.60 ± 103.08	236.39 ± 157.03	302.56 ± 70.79	<b>313.56 ± 52.17</b>
$T_1$	174.15 ± 170.30	253.56 ± 72.66	220.28 ± 118.61	<b>264.69 ± 61.63</b>	203.93 ± 167.83	241.45 ± 124.54	<b>241.60 ± 139.93</b>
$T_2$	135.16 ± 182.30	243.27 ± 89.86	<b>265.37 ± 80.60</b>	255.21 ± 90.61	84.10 ± 198.12	198.20 ± 151.64	<b>229.75 ± 166.87</b>

Table 2: Reward values gained by LRPG and baselines on continuous control tasks.

328 **Robustness Results: Continuous Control.** We studied the effectiveness of LRPG on continuous  
 329 control problems, and compared LR-PPO and LR-SAC to baselines for three different continuous  
 330 control environments on Gymnasium [Brockman et al., 2016]: *MountainCarContinuous*, *LunarLan-*  
 331 *derContinuous* and *BipedalWalker-v3*. Again, we trained 10 independent agents, and reported the  
 332 scores of the median agent. The results for the different noise kernels tested are presented in Table 5.

## 333 7 Discussion

334 **Experiments.** We applied LRPG on PPO, A2C and SAC algorithms, for a set of discrete and contin-  
 335 uous control environments. These environments are particularly sensitive to robustness problems;  
 336 the rewards are sparse, and applying a sub-optimal action at any step of the trajectory often leads to  
 337 terminal states with zero (or negative) reward. LRPG successfully induces lower robustness regrets in  
 338 the tested scenarios, and the use of  $K_D$  as an objective (even though we did not prove the convergence  
 339 of a gradient based method with such objective) yields a better compromise between robustness  
 340 and rewards. When compared to recent observational robustness methods, LRPG obtains similar  
 341 robustness results while *preserving the original guarantees of the chosen algorithm*<sup>7</sup>. However, the  
 342 improvements seem to be smaller for SAC, possibly due to the different nature of policy losses (SAC  
 343 uses a Q function as a loss).

344 **Shortcomings.** The motivation for LRPG comes from situations where, when deploying a model-free  
 345 controller in a dynamical system, we do not have a way of estimating the noise generation. There is  
 346 an alternative approach for robust RL, exploited in the reviewed literature, which consists in assuming  
 347 a disturbance structure (*e.g.* adversarial noise) and training directly to optimise the rewards in the  
 348 disturbed MDP. There is no clear answer on what approach is more rational, or more effective in  
 349 practice. The choice would depend on the problem at hand, the possible existence of an adversary,  
 350 the requirement (or lack thereof) for formal guarantees, *etc*. We cannot claim that our approach is  
 351 better in every way; we show through this work that LRPG is a useful approach for learning policies  
 352 in control problems where the noise sources are unknown and *we need to retain certain formal*  
 353 *guarantees* of the algorithms used. However, training against adversarial noise would possibly yield  
 354 higher robustness if tuned properly on many problems. An interesting direction would be to prove  
 355 preservation of guarantees for adversarial noise losses.

356 **Robustness, Complexity and Invariances.** Sections 2 and 3 discuss at large the structure, shape and  
 357 dependence of the maximally robust policy sets. These insights help derive optimisation objectives to  
 358 use in LRPG, but there is more to be said about how policy robustness is affected by the underlying  
 359 MDP properties. We hint at this in the proof of Corollary 3.5. More regular (*less complex* in entropy  
 360 terms, or more *symmetric*) reward functions (*e.g.*, reward functions with smaller variance across  
 361 the actions  $R(x, \cdot, y)$ ) seem to induce larger robust policy sets. In other words, for a fixed policy, a  
 362 *more complex* reward function yields larger robustness regrets as soon as any noise is introduced in  
 363 the system. This raises questions on how to use these principles to derive more robust policies in a  
 364 comprehensive way, but we leave these questions for future work.

<sup>7</sup>it even outperforms in some cases, although this is probably highly problem dependent, so we do not claim an improvement for every DOMDP

365 **References**

- 366 Mohammed Amin Abdullah, Hang Ren, Haitham Bou Ammar, Vladimir Milenkovic, Rui Luo,  
367 Mingtian Zhang, and Jun Wang. Wasserstein robust reinforcement learning. *arXiv preprint  
368 arXiv:1907.13196*, 2019.
- 369 Alekh Agarwal, Sham M Kakade, Jason D Lee, and Gaurav Mahajan. On the theory of policy  
370 gradient methods: Optimality, approximation, and distribution shift. *J. Mach. Learn. Res.*, 22(98):  
371 1–76, 2021.
- 372 Vahid Behzadan and Arslan Munir. Vulnerability of deep reinforcement learning to policy induction  
373 attacks. In *International Conference on Machine Learning and Data Mining in Pattern Recognition*,  
374 pages 262–275. Springer, 2017.
- 375 A. Ben-Israel and B. Mond. What is invexity? *The Journal of the Australian Mathematical Society.  
Series B. Applied Mathematics*, 28(1):1–9, 1986a.
- 377 Adi Ben-Israel and Bertram Mond. What is invexity? *The ANZIAM Journal*, 28(1):1–9, 1986b.
- 378 Dimitri Bertsekas. *Nonlinear Programming*. Athena Scientific, 1999.
- 379 Dimitri P Bertsekas. Nonlinear programming. *Journal of the Operational Research Society*, 48(3):  
380 334–334, 1997.
- 381 Jalaj Bhandari and Daniel Russo. Global optimality guarantees for policy gradient methods. *arXiv  
382 preprint arXiv:1906.01786*, 2019.
- 383 Shalabh Bhatnagar, Richard S. Sutton, Mohammad Ghavamzadeh, and Mark Lee. Natural actor–critic  
384 algorithms. *Automatica*, 45(11):2471–2482, November 2009. doi: 10.1016/j.automatica.2009.07.  
385 008.
- 386 Vivek S. Borkar. *Stochastic Approximation*. Hindustan Book Agency, 2008.
- 387 V.S. Borkar and K. Soumyanatha. An analog scheme for fixed point computation. i. theory. *IEEE  
388 Transactions on Circuits and Systems I: Fundamental Theory and Applications*, 44(4):351–355,  
389 1997. doi: 10.1109/81.563625.
- 390 Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and  
391 Wojciech Zaremba. Openai gym. *arXiv preprint arXiv:1606.01540*, 2016.
- 392 Maxime Chevalier-Boisvert, Lucas Willems, and Suman Pal. Minimalistic gridworld environment  
393 for openai gym. <https://github.com/maximecb/gym-minigrid>, 2018.
- 394 Esther Derman, Daniel Mankowitz, Timothy Mann, and Shie Mannor. A bayesian approach to robust  
395 reinforcement learning. In *Uncertainty in Artificial Intelligence*, pages 648–658. PMLR, 2020.
- 396 Michael Everett, Björn Lütjens, and Jonathan P How. Certifiable robustness to adversarial state  
397 uncertainty in deep reinforcement learning. *IEEE Transactions on Neural Networks and Learning  
398 Systems*, 2021.
- 399 Marc Fischer, Matthew Mirman, Steven Stalder, and Martin Vechev. Online robustness training for  
400 deep reinforcement learning. *arXiv preprint arXiv:1911.00887*, 2019.
- 401 Justin Fu, Katie Luo, and Sergey Levine. Learning robust rewards with adversarial inverse reinforce-  
402 ment learning. In *International Conference on Learning Representations*, 2018.
- 403 Adam Gleave, Michael Dennis, Cody Wild, Neel Kant, Sergey Levine, and Stuart Russell. Adver-  
404 sarial policies: Attacking deep reinforcement learning. In *International Conference on Learning  
405 Representations*, 2020.
- 406 Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy  
407 maximum entropy deep reinforcement learning with a stochastic actor. In *International conference  
408 on machine learning*, pages 1861–1870. PMLR, 2018.
- 409 Morgan A Hanson. On sufficiency of the kuhn-tucker conditions. *Journal of Mathematical Analysis  
410 and Applications*, 80(2):545–550, 1981.

- 411 Matthias Heger. Consideration of risk in reinforcement learning. In *Machine Learning Proceedings*  
 412 1994, pages 105–111. Elsevier, 1994.
- 413 Roger A Horn and Charles R Johnson. *Matrix analysis*. Cambridge university press, 2012.
- 414 Sandy Huang, Nicolas Papernot, Ian Goodfellow, Yan Duan, and Pieter Abbeel. Adversarial attacks  
 415 on neural network policies. *arXiv preprint arXiv:1702.02284*, 2017.
- 416 H Isermann. Linear lexicographic optimization. *Operations-Research-Spektrum*, 4(4):223–228, 1982.
- 417 Richard Klima, Daan Bloembergen, Michael Kaisers, and Karl Tuyls. Robust temporal difference  
 418 learning for critical domains. In *Proceedings of the 18th International Conference on Autonomous*  
 419 *Agents and MultiAgent Systems*, AAMAS ’19, page 350–358, Richland, SC, 2019. International  
 420 Foundation for Autonomous Agents and Multiagent Systems. ISBN 9781450363099.
- 421 Vijay R. Konda and John N. Tsitsiklis. Actor-critic algorithms. In S. A. Solla, T. K. Leen, and  
 422 K. Müller, editors, *Advances in Neural Information Processing Systems 12*, pages 1008–1014. MIT  
 423 Press, 2000.
- 424 Jernej Kos and Dawn Song. Delving into adversarial attacks on deep policies. *arXiv preprint*  
 425 *arXiv:1705.06452*, 2017.
- 426 Yen-Chen Lin, Zhang-Wei Hong, Yuan-Hong Liao, Meng-Li Shih, Ming-Yu Liu, and Min Sun.  
 427 Tactics of adversarial attack on deep reinforcement learning agents. In *Proceedings of the 26th*  
 428 *International Joint Conference on Artificial Intelligence*, pages 3756–3762, 2017.
- 429 Boyi Liu, Qi Cai, Zhuoran Yang, and Zhaoran Wang. Neural trust region/proximal policy optimization  
 430 attains globally optimal policy. In *Advances in Neural Information Processing Systems*, volume 32.  
 431 Curran Associates, Inc., 2019.
- 432 Ajay Mandlekar, Yuke Zhu, Animesh Garg, Li Fei-Fei, and Silvio Savarese. Adversarially robust  
 433 policy learning: Active construction of physically-plausible perturbations. In *2017 IEEE/RSJ*  
 434 *International Conference on Intelligent Robots and Systems (IROS)*, pages 3932–3939. IEEE,  
 435 2017.
- 436 Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, Alex Graves, Timothy Lillicrap, Tim  
 437 Harley, David Silver, and Koray Kavukcuoglu. Asynchronous methods for deep reinforcement  
 438 learning. In *International conference on machine learning*, pages 1928–1937. PMLR, 2016.
- 439 Janosch Moos, Kay Hansel, Hany Abdulsamad, Svenja Stark, Debora Clever, and Jan Peters. Robust  
 440 reinforcement learning: A review of foundations and recent advances. *Machine Learning and*  
 441 *Knowledge Extraction*, 4(1):276–315, 2022.
- 442 Jun Morimoto and Kenji Doya. Robust reinforcement learning. *Neural computation*, 17(2):335–359,  
 443 2005.
- 444 Andrew Y Ng, Daishi Harada, and Stuart Russell. Policy invariance under reward transformations:  
 445 Theory and application to reward shaping. In *Proc. of the Sixteenth International Conference on*  
 446 *Machine Learning*, 1999, 1999.
- 447 Arnab Nilim and Laurent El Ghaoui. Robust control of markov decision processes with uncertain  
 448 transition matrices. *Operations Research*, 53(5):780–798, 2005.
- 449 Xinlei Pan, Daniel Seita, Yang Gao, and John Canny. Risk averse robust adversarial reinforcement  
 450 learning. In *2019 International Conference on Robotics and Automation (ICRA)*, pages 8522–8528.  
 451 IEEE, 2019.
- 452 Santiago Paternain, Luiz F. O. Chamon, Miguel Calvo-Fullana, and Alejandro Ribeiro. Constrained  
 453 reinforcement learning has zero duality gap. In *Proceedings of the 33rd International Conference*  
 454 *on Neural Information Processing Systems*, pages 7553–7563, 2019.
- 455 Anay Pattanaik, Zhenyi Tang, Shuijing Liu, Gautham Bommannan, and Girish Chowdhary. Robust  
 456 deep reinforcement learning with adversarial attacks. *arXiv preprint arXiv:1712.03632*, 2017.

- 457 Anay Pattanaik, Zhenyi Tang, Shuijing Liu, Gautham Bommannan, and Girish Chowdhary. Robust  
 458 deep reinforcement learning with adversarial attacks. In *Proceedings of the 17th International*  
 459 *Conference on Autonomous Agents and MultiAgent Systems*, AAMAS '18, page 2040–2042,  
 460 Richland, SC, 2018. International Foundation for Autonomous Agents and Multiagent Systems.
- 461 Lerrel Pinto, James Davidson, Rahul Sukthankar, and Abhinav Gupta. Robust adversarial reinforce-  
 462 ment learning. In *International Conference on Machine Learning*, pages 2817–2826. PMLR,  
 463 2017.
- 464 Matteo Pirotta, Marcello Restelli, Alessio Pecorino, and Daniele Calandriello. Safe policy iteration.  
 465 In *International Conference on Machine Learning*, pages 307–315. PMLR, 2013.
- 466 Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah  
 467 Dormann. Stable-baselines3: Reliable reinforcement learning implementations. *Journal of Machine*  
 468 *Learning Research*, 22(268):1–8, 2021. URL <http://jmlr.org/papers/v22/20-1364.html>.
- 469 Mark J Rentmeesters, Wei K Tsai, and Kwei-Jay Lin. A theory of lexicographic multi-criteria  
 470 optimization. In *Proceedings of ICECCS'96: 2nd IEEE International Conference on Engineering*  
 471 *of Complex Computer Systems (held jointly with 6th CSESAW and 4th IEEE RTAW)*, pages 76–79.  
 472 IEEE, 1996.
- 473 Jay K Satia and Roy E Lave Jr. Markovian decision processes with uncertain transition probabilities.  
 474 *Operations Research*, 21(3):728–740, 1973.
- 475 John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy  
 476 optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- 477 Joar Skalse, Matthew Farrugia-Roberts, Stuart Russell, Alessandro Abate, and Adam Gleave. In-  
 478 variance in policy optimisation and partial identifiability in reward learning. *arXiv preprint*  
 479 *arXiv:2203.07475*, 2022a.
- 480 Joar Skalse, Lewis Hammond, Charlie Griffin, and Alessandro Abate. Lexicographic multi-objective  
 481 reinforcement learning. In *Proceedings of the Thirty-First International Joint Conference on*  
 482 *Artificial Intelligence*, pages 3430–3436, jul 2022b. doi: 10.24963/ijcai.2022/476.
- 483 Morton Slater. Lagrange multipliers revisited. Cowles Commission Discussion Paper No. 403, 1950.
- 484 Matthijs TJ Spaan. Partially observable markov decision processes. In *Reinforcement Learning*,  
 485 pages 387–414. Springer, 2012.
- 486 Matthijs TJ Spaan and N Vlassis. A point-based pomdp algorithm for robot planning. In *IEEE Inter-  
 487 national Conference on Robotics and Automation, 2004. Proceedings. ICRA'04. 2004*, volume 3,  
 488 pages 2399–2404. IEEE, 2004.
- 489 Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
- 490 Richard S. Sutton, David McAllester, Satinder Singh, and Yishay Mansour. Policy gradient methods  
 491 for reinforcement learning with function approximation. In *Proceedings of the 12th International*  
 492 *Conference on Neural Information Processing Systems*, NIPS'99, pages 1057–1063, Cambridge,  
 493 MA, USA, 1999. MIT Press.
- 494 Kai Liang Tan, Yasaman Esfandiari, Xian Yeow Lee, Soumik Sarkar, et al. Robustifying reinforce-  
 495 ment learning agents via action space adversarial training. In *2020 American control conference (ACC)*,  
 496 pages 3959–3964. IEEE, 2020.
- 497 Chen Tessler, Yonathan Efroni, and Shie Mannor. Action robust reinforcement learning and applica-  
 498 tions in continuous control. In *International Conference on Machine Learning*, pages 6215–6224.  
 499 PMLR, 2019.
- 500 Elise van der Pol, Thomas Kipf, Frans A. Oliehoek, and Max Welling. Plannable approximations  
 501 to mdp homomorphisms: Equivariance under actions. In *Proceedings of the 19th International*  
 502 *Conference on Autonomous Agents and MultiAgent Systems*, AAMAS '20, page 1431–1439,  
 503 Richland, SC, 2020. International Foundation for Autonomous Agents and Multiagent Systems.  
 504 ISBN 9781450375184.

- 505 Bart PG Van Parys, Daniel Kuhn, Paul J Goulart, and Manfred Morari. Distributionally robust control  
506 of constrained stochastic systems. *IEEE Transactions on Automatic Control*, 61(2):430–442, 2015.
- 507 Wolfram Wiesemann, Daniel Kuhn, and Melvyn Sim. Distributionally robust convex optimization.  
508 *Operations Research*, 62(6):1358–1376, 2014.
- 509 Huan Xu and Shie Mannor. The robustness-performance tradeoff in markov decision processes.  
510 *Advances in Neural Information Processing Systems*, 19, 2006.
- 511 Huan Zhang, Hongge Chen, Chaowei Xiao, Bo Li, Mingyan Liu, Duane Boning, and Cho-Jui  
512 Hsieh. Robust deep reinforcement learning against adversarial perturbations on state observations.  
513 *Advances in Neural Information Processing Systems*, 33:21024–21037, 2020.
- 514 Huan Zhang, Hongge Chen, Duane Boning, and Cho-Jui Hsieh. Robust reinforcement learning  
515 on state observations with learned optimal adversary. In *International Conference on Learning  
Representation (ICLR)*, 2021.
- 517 Shangtong Zhang. Modularized implementation of deep rl algorithms in pytorch. <https://github.com/ShangtongZhang/DeepRL>, 2018.

519 **A Examples and Further Considerations**

520 We provide here two examples to show how we can obtain limit scenarios  $\Pi_0 = \Pi$  (any policy is  
 521 maximally robust) or  $\Pi_0 = \Pi_T$  (Example 1), and how for some MDPs the third inclusion in Theorem  
 522 3.4 is strict (Example 2).

523 **Example 1** Consider the simple MDP in Figure 2. First, consider the reward function  $R_1(x_1, \cdot, \cdot) =$   
 524  $10$ ,  $R_1(x_2, \cdot, \cdot) = 0$ . This produces a “dummy” MDP where all policies have the same reward sum.  
 525 Then,  $\forall T, \pi, V^{\langle \pi, T \rangle} = V^\pi$ , and therefore we have  $\Pi_D = \Pi_0 = \Pi$ .

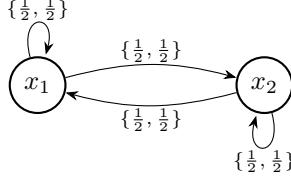


Figure 2: Example MDP. Values in brackets represent  $\{P(\cdot, u_1, \cdot), P(\cdot, u_2, \cdot)\}$ .

526 Now, consider the reward function  $R_2(x_1, u_1, \cdot) = 10$ ,  $R_2(\cdot, \cdot, \cdot) = 0$  elsewhere. Take a non-  
 527 constant policy  $\pi$ , i.e.,  $\pi(x_1) \neq \pi(x_2)$ . In the example DOMDP (assuming the initial state is drawn  
 528 uniformly from  $X_0 = \{x_1, x_2\}$ ) one can show that at any time in the trajectory, there is a stationary  
 529 probability  $\Pr\{x_t = x_1\} = \frac{1}{2}$ . Let us abuse notation and write  $\pi(x_i) = (\pi(x_i, u_1) \ \pi(x_i, u_2))^T$   
 530 and  $R(x_i) = (R(x_i, u_1, \cdot) \ R(x_i, u_2, \cdot))^T$ . For the given reward structure we have  $R(x_2) =$   
 531  $(0 \ 0)^T$ , and therefore:

$$J(\pi) = E_{x_0 \sim \mu_0} \left[ \sum_{t=0}^{\infty} \gamma^t R_t \right] = \frac{1}{2} \langle R(x_1), \pi(x_1) \rangle \frac{\gamma}{1-\gamma}. \quad (7)$$

Since the transitions of the MDP are independent of the actions, following the same principle as in (7):  $J(\langle \pi, T \rangle) = \frac{1}{2} \langle R(x_1), \langle \cdot, T \rangle(\pi)(x_1) \rangle \frac{\gamma}{1-\gamma}$ . For any noise map  $\langle \cdot, T \rangle \neq \text{Id}$ , for the two-state policy it holds that  $\pi \notin \Pi_T \implies \langle \pi, T \rangle \neq \pi$ . Therefore  $\langle \pi, T \rangle(x_1) \neq \pi(x_1)$  and:

$$J(\pi) - J(\langle \pi, T \rangle) = \frac{5\gamma}{1-\gamma} \cdot (\pi(x_1, 1) - \langle \pi, T \rangle(x_1, 1)) \neq 0,$$

532 which implies that  $\pi \notin \Pi_0$ .

533 **Example 2** Consider the same MDP in Figure 2 with reward function  $R(x_1, u_1, \cdot) = R(x_2, u_1, \cdot) =$   
 534  $10$ , and a reward of zero for all other transitions. Take a policy  $\pi(x_1) = (1 \ 0)$ ,  $\pi(x_2) = (0 \ 1)$ . The  
 535 policy yields a reward of  $10$  in state  $x_1$  and a reward of  $0$  in state  $x_2$ . Again we assume the initial  
 536 state is drawn uniformly from  $X_0 = \{x_1, x_2\}$ . Then, observe:

$$J(\pi) = E_{x_0 \sim \mu_0} \left[ \sum_{t=0}^{\infty} \gamma^t R_t \right] = \frac{1}{2} \langle R(x_1), \pi(x_1) \rangle \frac{\gamma}{1-\gamma} = \frac{1}{2} \frac{10\gamma}{1-\gamma} = \frac{5\gamma}{1-\gamma}.$$

537 Define now noise map  $T(\cdot | x_1) = (\frac{1}{2} \ \frac{1}{2})$  and  $T(\cdot | x_2) = (\frac{1}{2} \ \frac{1}{2})$ . Observe this noise map  
 538 yields a policy with non-zero disadvantage,  $D^\pi(x_1, T) = \frac{5\gamma}{1-\gamma} - (\frac{5\gamma}{1-\gamma} - 2.5) = 2.5$  and similarly  
 539  $D^\pi(x_2, T) = -2.5$ , therefore  $\pi \notin \Pi_D$ . However, the policy is *maximally robust*:

$$J(\langle \pi, T \rangle) = \frac{1}{2} \langle R(x_1), \langle \pi, T \rangle(x_1) \rangle \frac{\gamma}{1-\gamma} + \frac{1}{2} \langle R(x_2), \langle \pi, T \rangle(x_2) \rangle \frac{\gamma}{1-\gamma} = \frac{1}{2} \frac{\gamma}{1-\gamma} (5+5) = \frac{5\gamma}{1-\gamma}. \quad (8)$$

540 Therefore,  $\pi \in \Pi_0$ .

541 **B Theoretical Results**

542 **B.1 Auxiliary Results**

543 **Theorem B.1** (Stochastic Approximation with Non-Expansive Operator). *Let  $\{\xi_t\}$  be a random sequence with  $\xi_t \in \mathbb{R}^n$  defined by the iteration:*

$$\xi_{t+1} = \xi_t + \alpha_t(F(\xi_t) - \xi_t + M_{t+1}),$$

545 where:

- 546 1. The step sizes  $\alpha_t$  satisfy Assumption 2.
- 547 2.  $F : \mathbb{R}^n \mapsto \mathbb{R}^n$  is a  $\|\cdot\|_\infty$  non-expansive map. That is, for any  $\xi_1, \xi_2 \in \mathbb{R}^n$ ,  $\|F(\xi_1) - F(\xi_2)\|_\infty \leq \|\xi_1 - \xi_2\|_\infty$ .
- 549 3.  $\{M_t\}$  is a martingale difference sequence with respect to the increasing family of  $\sigma$ -fields  $550 \mathcal{F}_t := \sigma(\xi_0, M_0, \xi_1, M_1, \dots, \xi_t, M_t)$ .

551 Then, the sequence  $\xi_t \rightarrow \xi^*$  almost surely where  $\xi^*$  is a fixed point such that  $F(\xi^*) = \xi^*$ .

552 *Proof.* See Borkar and Soumyanatha [1997].  $\square$

553 **Theorem B.2** (PB-LRL Convergence). *Let  $\mathcal{M}$  be a multi-objective MDP with objectives  $K_i$ ,  $i \in 554 \{1, \dots, m\}$  of the same form. Assume a policy  $\pi$  is twice differentiable in parameters  $\theta$ , and if using a 555 critic  $V_i$  assume it is continuously differentiable on  $w_i$ . Suppose that if PB-LRL is run for  $T$  steps, 556 there exists some limit point  $w_i^*(\theta)$  when  $\theta$  is held fixed under conditions  $\mathcal{C}$  on  $\mathcal{M}$ ,  $\pi$  and  $V_i$ . If 557  $\lim_{T \rightarrow \infty} \mathbb{E}_t[\theta] \in \Theta_1^\epsilon$  for  $m = 1$ , then for any  $m \in \mathbb{N}$  we have  $\lim_{T \rightarrow \infty} \mathbb{E}_t[\theta] \in \Theta_m^\epsilon$  where  $\epsilon$  depends 558 on the representational power of the parameterisations of  $\pi$ ,  $V_i$ .*

559 *Proof Sketch.* We refer the interested reader to Skalse et al. [2022b] for a full proof, and here attempt 560 to provide the intuition behind the result in the form of a proof sketch.

561 Let us begin by briefly recalling the general problem statement: we wish to take a multi-objective 562 MDP  $\mathcal{M}$  with  $m$  objectives, and obtain a lexicographically optimal policy (one that optimises the 563 first objective, and then subject to this optimises the second objective, and so on). More precisely, 564 for a policy  $\pi$  parameterised by  $\theta$ , we say that  $\pi$  is (globally) *lexicographically  $\epsilon$ -optimal* if  $\theta \in \Theta_m^\epsilon$ , 565 where  $\Theta_0^\epsilon = \Theta$  is the set of all policies in  $\mathcal{M}$ ,  $\Theta_{i+1}^\epsilon := \{\theta \in \Theta_i^\epsilon \mid \max_{\theta' \in \Theta_i^\epsilon} K_i(\theta') - K_i(\theta) \leq \epsilon_i\}$ , 566 and  $\mathbb{R}^{m-1} \ni \epsilon \succcurlyeq 0$ .<sup>8</sup>

567 The basic idea behind policy-based lexicographic reinforcement learning (PB-LRL) is to use a multi- 568 timescale approach to first optimise  $\theta$  using  $K_1$ , then at a slower timescale optimise  $\theta$  using  $K_2$  while 569 adding the condition that the loss with respect to  $K_1$  remains bounded by its current value, and so on. 570 This sequence of constrained optimisations problems can be solved using a Lagrangian relaxation 571 [Bertsekas, 1999], either in series or – via a judicious choice of learning rates – simultaneously, by 572 exploiting a separation in timescales [Borkar, 2008]. In the simultaneous case, the parameters of the 573 critic  $w_i$  (if using an actor-critic algorithm, if not this part of the argument may be safely ignored) 574 for each objective are updated on the fastest timescale, then the parameters  $\theta$ , and finally (i.e., most 575 slowly) the Lagrange multipliers for each of the remaining constraints.

576 The proof proceeds via induction on the number of objectives, using a standard stochastic approxi- 577 mation argument [Borkar, 2008]. In particular, due to the learning rates chosen, we may consider 578 those more slowly updated parameters fixed for the purposes of analysing the convergence of the 579 more quickly updated parameters. In the base case where  $m = 1$ , we have (by assumption) that 580  $\lim_{T \rightarrow \infty} \mathbb{E}_t[\theta] \in \Theta_1^\epsilon$ . This is simply the standard (non-lexicographic) RL setting. Before continuing 581 to the inductive step, Skalse et al. [2022b] observe that because gradient descent on  $K_1$  converges to 582 globally optimal stationary point when  $m = 1$  then  $K_1$  must be globally *invex* (where the opposite 583 implication is also true) [Ben-Israel and Mond, 1986a].<sup>9</sup>

<sup>8</sup>The proof in Skalse et al. [2022b] also considers *local* lexicographic optima, though for the sake of simplicity, we do not do so here.

<sup>9</sup>A differentiable function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  is (globally) *invex* if and only if there exists a function  $g : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}^n$  such that  $f(x_1) - f(x_2) \geq g(x_1, x_2)^\top \nabla f(x_2)$  for all  $x_1, x_2 \in \mathbb{R}^n$  [Hanson, 1981].

584 The reason this observation is useful is that because each of the objectives  $K_i$  shares the same  
 585 functional form, they are all invex, and furthermore, invexity is conserved under linear combinations  
 586 and the addition of scalars, meaning that the Lagrangian formed in the relaxation of each constrained  
 587 optimisation problem is also invex. As a result, if we assume that  $\lim_{T \rightarrow \infty} \mathbb{E}_t[\theta] \in \Theta_i^\epsilon$  as our  
 588 inductive hypothesis, then the stationary point of the Lagrangian for optimising objective  $K_{i+1}$  is  
 589 a global optimum, given the constraints that it does not worsen performance on  $K_1, \dots, K_i$ . Via  
 590 Slater's condition [Slater, 1950] and standard saddle-point arguments [Bertsekas, 1999, Paternain  
 591 et al., 2019], we therefore have that  $\lim_{T \rightarrow \infty} \mathbb{E}_t[\theta] \in \Theta_{i+1}^\epsilon$ , completing the inductive step, and thus  
 592 the overall inductive argument.

593 This concludes the proof that  $\lim_{T \rightarrow \infty} \mathbb{E}_t[\theta] \in \Theta_m^\epsilon$ . We refer the reader to Skalse et al. [2022b] for a  
 594 discussion of the error  $\epsilon$ , but intuitively it corresponds to a combination of the representational power  
 595 of  $\theta$ , the critic parameters  $w_i$  (if used), and the duality gap due to the Lagrangian relaxation [Paternain  
 596 et al., 2019]. In cases where the representational power of the various parameters is sufficiently high,  
 597 then it can be shown that  $\epsilon = 0$ .  $\square$

598 **Lemma B.3.** *Let  $\pi_\theta$  be a fully-parameterised policy in a DOMDP, and  $\alpha_t$  a learning rate satisfying  
 599 Assumption 2. Consider the following approximated gradient for objective  $K_{\tilde{T}}(\pi)$  and sampled point  
 600  $x \in X$ :*

$$\nabla_\theta \hat{K}_{\tilde{T}}(\theta)(x) = (\pi_\theta(x) - \pi_\theta(y)) \nabla_\theta \pi_\theta(x), \quad y \sim \tilde{T}(\cdot | x). \quad (9)$$

601 Then, the following iteration with  $x \in X$  and some initial  $\theta_0$ ,

$$\theta_{t+1} = \theta_t - \alpha_t \nabla_\theta \hat{K}_{\tilde{T}}(\theta_t) \quad (10)$$

602 yields  $\theta \rightarrow \tilde{\theta}$  almost surely where  $\tilde{\theta}$  satisfies  $K_{\tilde{T}}(\tilde{\theta}) = 0$ .

603 *Lemma B.3.* We make use of standard results on stochastic approximation with non-expansive  
 604 operators (specifically, Theorem B.1 in the appendix) Borkar and Soumyanatha [1997]. First, observe  
 605 that for a fully parameterised policy, one can assume to have a tabular representation such that  
 606  $\pi_\theta(x, u) = \theta_{xu}$ , and  $\nabla_\theta \pi_\theta(x) \equiv \text{Id}$ . We can then write the stochastic gradient descent problem in  
 607 terms of the policy. Let  $y \sim \tilde{T}(\cdot | x)$ . Then:

$$\begin{aligned} \pi_{t+1}(x) &= \pi_t(x) - \alpha_t (\pi_t(x) - \pi_t(y)) = \\ &= \pi_t(x) - \alpha_t \left( \pi_t(x) - \langle \pi_t, \tilde{T} \rangle(x) - (\pi_t(y) - \langle \pi_t, \tilde{T} \rangle(x)) \right). \end{aligned}$$

608 We now need to verify that the necessary conditions for applying Theorem B.1 hold. First,  $\alpha_t$  satisfies  
 609 Assumption 2. Second, making use of the property  $\|\tilde{T}\|_\infty = 1$  for any row-stochastic matrix  $\tilde{T}$ , for  
 610 any two policies  $\pi_1, \pi_2 \in \Pi$ :

$$\|\langle \pi_1, \tilde{T} \rangle - \langle \pi_2, \tilde{T} \rangle\|_\infty = \|\tilde{T}\pi_1 - \tilde{T}\pi_2\|_\infty = \|\tilde{T}(\pi_1 - \pi_2)\|_\infty \leq \|\tilde{T}\|_\infty \|\pi_1 - \pi_2\|_\infty = \|\pi_1 - \pi_2\|_\infty.$$

611 Therefore, the operator  $\langle \cdot, \tilde{T} \rangle$  is non-expansive with respect to the sup-norm. For the final condition,  
 612 we have

$$\mathbb{E}_{y \sim \tilde{T}(\cdot | x)} [\pi_t(y) - \langle \pi_t, \tilde{T} \rangle(x) | \pi_t, \tilde{T}] = \sum_{y \in X} \tilde{T}(y | x) \pi_t(y) - \langle \pi_t, \tilde{T} \rangle(x) = 0.$$

613 Therefore, the difference  $\pi_t(y) - \langle \pi_t, \tilde{T} \rangle(x)$  is a martingale difference for all  $x$ . One can then apply  
 614 Theorem B.1 with  $\xi_t(x) \equiv \pi_t(x)$ ,  $F(\cdot) \equiv \langle \cdot, \tilde{T} \rangle$  and  $M_{t+1} \equiv \pi_t(y) - \langle \pi_t, \tilde{T} \rangle(x)$  to conclude that  
 615  $\pi_t(x) \rightarrow \tilde{\pi}(x)$  almost surely. Finally from assumption 2.2, for any policy all states  $x \in X$  are visited  
 616 infinitely often, therefore  $\pi_t(x) \rightarrow \tilde{\pi}(x) \forall x \in X \implies \pi_t \rightarrow \tilde{\pi}$  and  $\tilde{\pi}$  satisfies  $\langle \tilde{\pi}, \tilde{T} \rangle = \tilde{\pi}$ , and  
 617  $K_{\tilde{T}}(\tilde{\pi}) = 0$ .  $\square$

## 618 B.2 Proofs

619 We now present the proofs for the statements through the work.

620 **Proposition 3.2.** If a policy  $\pi \in \Pi$  is a fixed point of the operator  $\langle \cdot, T \rangle$ , then it holds that  $\langle \pi, T \rangle = \pi$ .  
 621 Therefore, one can compute the robustness of the policy  $\pi$  to obtain  $\rho(\pi, T) = J(\pi) - J(\langle \pi, T \rangle) =$   
 622  $J(\pi) - J(\pi) = 0 \implies \pi \in \Pi_0$ . Therefore,  $\Pi_T \subseteq \Pi_0$ .

623 For a discrete state and action spaces, the space of stochastic kernels  $\mathcal{K} : X \mapsto \Delta(X)$  is equivalent  
624 to the space of row-stochastic  $|X| \times |X|$  matrices, therefore one can write  $T(y \mid x) \equiv T_{xy}$  as the  
625  $xy$ -th entry of the matrix  $T$ . Then, the representation of a constant policy as an  $X \times U$  matrix can be  
626 written as  $\bar{\pi} = \mathbf{1}_{|X|} v^\top$ , where  $\mathbf{1}_{|X|}$  where  $v \in \Delta(U)$  is any probability distribution over the action  
627 space. Observe that, applying the operator  $\langle \pi, T \rangle$  to a constant policy yields:

$$\langle \bar{\pi}, T \rangle = T \mathbf{1}_{|X|} v^\top. \quad (11)$$

628 By the Perron-Frobenius Theorem [Horn and Johnson, 2012], since  $T$  is row-stochastic it has at least  
629 one eigenvalue  $\text{eig}(T) = 1$ , and this admits a (strictly positive) eigenvector  $T \mathbf{1}_{|X|} = \mathbf{1}_{|X|}$ . Therefore,  
630 substituting this in (11):

$$\langle \bar{\pi}, T \rangle = T \mathbf{1}_{|X|} v^\top = \mathbf{1}_{|X|} v^\top = \bar{\pi} \implies \bar{\Pi} \subseteq \Pi_T.$$

631  $\square$

632 *Proposition 3.3.* Recall the definition in (2) and that the noise disadvantage function of a policy  $\pi$  is  
633 given by (4). We want to show that  $D^\pi(x, T) = 0 \implies \rho(\pi, T) = 0$ . Taking  $D^\pi(x, T) = 0$  one  
634 has a policy that produces an disadvantage of zero when noise kernel  $T$  is applied. Then,

$$D^\pi(x, T) = 0 \implies \mathbb{E}_{u \sim \langle \pi, T \rangle(x)}[Q^\pi(x, u)] = V^\pi(x) \quad \forall x \in X. \quad (12)$$

Now define the value of the disturbed policy

$$V^{\langle \pi, T \rangle}(x_0) := \mathbb{E}_{\substack{u_k \sim \langle \pi, T \rangle(x_k), \\ x_{k+1} \sim P(\cdot|x_k, u_k)}} \left[ \sum_{k=0}^{\infty} \gamma^k r(x_k, u_k) \right],$$

635 and take:

$$V^{\langle \pi, T \rangle}(x) = \mathbb{E}_{\substack{u \sim \langle \pi, T \rangle(x), \\ y \sim P(\cdot|x, u)}} \left[ r(x, u, y) + \gamma V^{\langle \pi, T \rangle}(y) \right].$$

636 We will now show that  $V^\pi(x) = V^{\langle \pi, T \rangle}(x)$ , for all  $x \in X$ . Observe, from (12) using  $V^\pi(x) =$   
637  $\mathbb{E}_{u \sim \langle \pi, T \rangle(x)}[Q^\pi(x, u)]$ , we have  $\forall x \in X$ :

$$\begin{aligned} V^\pi(x) - V^{\langle \pi, T \rangle}(x) &= \mathbb{E}_{u \sim \langle \pi, T \rangle(x)}[Q^\pi(x, u)] - \mathbb{E}_{\substack{u \sim \langle \pi, T \rangle(x), \\ y \sim P(\cdot|x, u)}} \left[ r(x, u, y) + \gamma V^{\langle \pi, T \rangle}(y) \right] \\ &= \mathbb{E}_{\substack{u \sim \langle \pi, T \rangle(x), \\ y \sim P(\cdot|x, u)}} \left[ r(x, u, y) + \gamma V^\pi(y) - r(x, u, y) - \gamma V^{\langle \pi, T \rangle}(y) \right] \quad (13) \\ &= \gamma \mathbb{E}_{y \sim P(\cdot|x, u)} \left[ V^\pi(y) - V^{\langle \pi, T \rangle}(y) \right]. \end{aligned}$$

638 Now, taking the sup norm at both sides of (13) we get

$$\|V^\pi(x) - V^{\langle \pi, T \rangle}(x)\|_\infty = \gamma \left\| \mathbb{E}_{y \sim P(\cdot|x, u)} \left[ V^\pi(y) - V^{\langle \pi, T \rangle}(y) \right] \right\|_\infty. \quad (14)$$

639 Observe that for the right hand side of (14), we have  $\|\mathbb{E}_{y \sim P(\cdot|x, u)} [V^\pi(y) - V^{\langle \pi, T \rangle}(y)]\|_\infty \leq$   
640  $\|V^\pi(x) - V^{\langle \pi, T \rangle}(x)\|_\infty$ . Therefore, since  $\gamma < 1$ ,

$$\|V^\pi(x) - V^{\langle \pi, T \rangle}(x)\|_\infty \leq \gamma \|V^\pi(x) - V^{\langle \pi, T \rangle}(x)\|_\infty \implies \|V^\pi(x) - V^{\langle \pi, T \rangle}(x)\|_\infty = 0. \quad (15)$$

641 Finally,  $\|V^\pi(x) - V^{\langle \pi, T \rangle}(x)\|_\infty = 0 \implies V^\pi(x) - V^{\langle \pi, T \rangle}(x) = 0 \quad \forall x \in X$ , and  $V^\pi(x) -$   
642  $V^{\langle \pi, T \rangle}(x) = 0 \quad \forall x \in X \implies J(\pi) = J(\langle \pi, T \rangle) \implies \rho(\pi, T) = 0$ .  $\square$

643 *Inclusion Theorem 3.4.* Combining Proposition 3.2 and Proposition 3.3, we simply need to show  
644 that  $\Pi_T \subset \Pi_D$ . Take  $\pi$  to be a fixed point of  $\langle \pi, T \rangle$ . Then  $\langle \pi, T \rangle = \pi$ , and from the definition in (4):

$$\begin{aligned} D^\pi(x, T) &= V^\pi(x) - \mathbb{E}_{u \sim \langle \pi, T \rangle(x, \cdot)}[Q^\pi(x, u)] \\ &= V^\pi(x) - \mathbb{E}_{u \sim \pi(x, \cdot)}[Q^\pi(x, u)] \\ &= V^\pi(x) - V^\pi(x) = 0. \end{aligned}$$

645 Therefore,  $\pi \in \Pi_D$ , which completes the sequence of inclusions.

646 To show convexity of  $\bar{\Pi}$ ,  $\Pi_T$ , first for a constant policy  $\bar{\pi} \in \bar{\Pi}$ , recall that we can write  $\bar{\pi} = \mathbf{1}v^\top$ ,  
 647 where  $v \in \Delta(U)$  is any probability distribution over the action space. Now take  $\bar{\pi}_1, \bar{\pi}_2 \in \bar{\Pi}$ . For any  
 648  $\alpha \in [0, 1]$ ,  $\alpha\bar{\pi}_1 + (1 - \alpha)\bar{\pi}_2 = \alpha\mathbf{1}v_1^\top + (1 - \alpha)\mathbf{1}v_2^\top = \mathbf{1}(\alpha v_1 + (1 - \alpha)v_2)^\top \in \bar{\Pi}$ .

649 At last, for the set  $\Pi_T$ , assume there exist two different policies  $\pi_1, \pi_2$  both fixed points of  $\langle \cdot, T \rangle$ .  
 650 Then, for any  $\alpha \in [0, 1]$ ,  $\langle (\alpha\pi_1 + (1 - \alpha)\pi_2), T \rangle = \alpha T\pi_1 + (1 - \alpha)T\pi_2 = \alpha\pi_1 + (1 - \alpha)\pi_2$ .  
 651 Therefore, any affine combination of fixed points is also a fixed point.  $\square$

652 *Corollary 3.5.* For statement (i), let  $\bar{R}(\cdot, \cdot, \cdot) = c$  for some constant  $c \in \mathbb{R}$ . Then,  $J(\pi) =$   
 653  $\mathbb{E}_{x_0 \sim \mu_0} [\sum_t \gamma^t \bar{r}_t \mid \pi] = \frac{c\gamma}{1-\gamma}$ , which does not depend on the policy  $\pi$ . For any noise kernel  $T$   
 654 and policy  $\pi$ ,  $J(\pi) - J(\langle \pi, T \rangle) = 0 \implies \pi \in \Pi_0$ .

For statement (ii) assume  $\exists \pi \in \Pi_0 : \pi \notin \Pi_T$ . Then,  $\exists x^* \in X$  and  $u^* \in U$  such that  $\pi(x^*, u^*) \neq$   
 $\langle \pi, T \rangle(x^*, u^*)$ . Let:

$$\underline{R}(x, u, x') := \begin{cases} c & \text{if } x = x^* \text{ and } u = u^* \\ 0 & \text{otherwise} \end{cases}.$$

655 Then,  $\mathbb{E}[R(x, \pi(x), x')] < \mathbb{E}[R(x, \langle \pi, T \rangle(x), x')]$  and since the MDP is ergodic  $x$  is visited infinitely  
 656 often and  $J(\pi) - J(\langle \pi, T \rangle) > 0 \implies \pi \notin \Pi_0$ , which contradicts the assumption. Therefore,  
 657  $\Pi_0 \setminus \Pi_T = \emptyset \implies \Pi_0 = \Pi_T$ .  $\square$

658 *Theorem 4.2.* We apply the results from Skalse et al. [2022b] in Theorem B.2. Essentially, Skalse  
 659 et al. [2022b] prove that for a policy gradient algorithm to lexicographically optimise a policy for  
 660 multiple objectives, it is a sufficient condition that the stochastic gradient descent algorithm finds  
 661 optimal parameters for each of the objectives independently. From Lemma B.3 we know that a policy  
 662 gradient algorithm using the gradient estimate in (9) converges to a maximally robust policy, i.e.  
 663 a set of parameters  $\theta' = \operatorname{argmax}_\theta K_{\tilde{T}}$ . Additionally, by assumption, the chosen algorithm for  $K_1$   
 664 converges to an optimal point  $\theta^*$ . While the two objective functions are not of the same form – as  
 665 in Skalse et al. [2022b] – the fact they are both invex [Ben-Israel and Mond, 1986b] either locally  
 666 or globally depending on the form of  $K_1$ , implies that  $\hat{K}$  is also invex and hence that the stationary  
 667 point  $\theta^e$  computed by LRPG satisfies equation 6.  $\square$

668 *Corollary 5.1.* The proof follows by the inclusion results in Theorem 3.4. If  $\Pi_{\tilde{T}} = \bar{\Pi}$ , then  $\Pi_{\tilde{T}} \subseteq \Pi_T$   
 669 for any other  $T$ . Then, the distance from  $\pi$  to the set  $\Pi_T$  is at most the distance to  $\Pi_{\tilde{T}}$ .  $\square$

### 670 B.3 On Adversarial Disturbances and other Noise Kernels

671 A problem that remains open after this work is what constitutes an appropriate choice of  $\tilde{T}$ , and what  
 672 can we expect by restricting a particular class of  $\tilde{T}$ . We first discuss adversarial examples, and then  
 673 general considerations on  $\tilde{T}$  versus  $T$ .

**Adversarial Noise** As mentioned in the introduction, much of the previous work focuses on  
 adversarial disturbances. We did not directly address this in the results of this work since our  
 motivation lies in the scenarios where the disturbance is not adversarial and is unknown. However,  
 following the results of Section 3, we are able to reason about adversarial disturbances. Consider an  
 adversarial map  $T_{adv}$  to be

$$\langle \pi, T_{adv} \rangle(x) = \pi(y), \quad y \in \operatorname{argmax}_{y \in X_{ad}(x)} d(\pi(x), \pi(y)),$$

674 with  $X_{ad}(x) \subseteq X$  being a set of admissible disturbance states for  $x$ , and  $d(\cdot, \cdot)$  is a distance measure  
 675 between distributions (e.g. 2-norm).

676 **Proposition B.4.** Constant policies are a fixed point of  $T_{adv}$ , and are the only fixed points if for all  
 677 pairs  $x_0, x_k$  there exists a sequence  $\{x_0, \dots, x_k\} \subseteq X$  such that  $x_i \in X_{ad}(x_i)$ .

678 *Proof.* First, it is straight-forward that if  $\bar{\pi} \in \bar{\Pi} \implies \langle \bar{\pi}, T_{adv} \rangle(x) = \bar{\pi}(x)$ . To show they are the  
 679 only fixed points, assume that there is a non-constant policy  $\pi'$  that is a fixed point of  $T_{ad}$ . Then,  
 680 there exists  $x, z$  such that  $\pi'(x) \neq \pi'(z)$ . However, by assumption, we can construct a sequence  
 681  $\{x, \dots, z\} \subseteq X$  that connects  $x$  and  $z$  and every state in the sequence is in the admissible set of  
 682 the previous one. Assume without loss of generality that this sequence is  $\{x, y, z\}$ . Then, if  $\pi'$  is

683 a fixed point,  $\langle \pi', T_{adv} \rangle(x) = \pi'(x)$ ,  $\langle \pi', T_{adv} \rangle(y) = \pi'(y)$  and  $\langle \pi', T_{adv} \rangle(z) = \pi'(z)$ . However,  
684  $\pi'(x) \neq \pi'(z)$ , so either  $\pi'(x) \neq \pi'(y) \implies d(\pi'(x), \pi'(y)) \neq 0$  or  $\pi'(y) \neq \pi'(z) \implies$   
685  $d(\pi'(y), \pi'(z)) \neq 0$ , therefore  $\pi'$  cannot be a fixed point of  $T_{adv}$ .  $\square$

686 The main difference between an adversarial operator and the random noise considered throughout  
687 this work is that  $T_{adv}$  is *not a linear operator*, and additionally, it is time varying (since the policy is  
688 being modified at every time step of the PG algorithm). Therefore, including it as a LRPG objective  
689 would invalidate the assumptions required for LRPG to retain formal guarantees of the original PG  
690 algorithm used, and it is not guaranteed that the resulting policy gradient algorithm would converge.

691 **C Experiment Methodology**

692 We use in the experiments well-tested implementations of A2C, PPO and SAC from Stable Baselines  
 693 3 [Raffin et al., 2021] to include the computation of the lexicographic parameters in (1).

694 **LRPG Parameters.** The LRL parameters are initialised in all cases as  $\beta_0^1 = 2$ ,  $\beta_0^2 = 1$ ,  $\lambda = 0$   
 695 and  $\eta = 0.001$ . The LRL tolerance is set to  $\epsilon_t = 0.99\hat{k}_1$  to ensure we never deviate too much  
 696 from the original objective, since the environments have very sparse rewards. We use a first order  
 697 approximation to compute the LRL weights from the original LMORL implementation.

698 **C.1 Discrete Control**

The discrete control environments used can be seen in Figure 3. Since all the environments use a

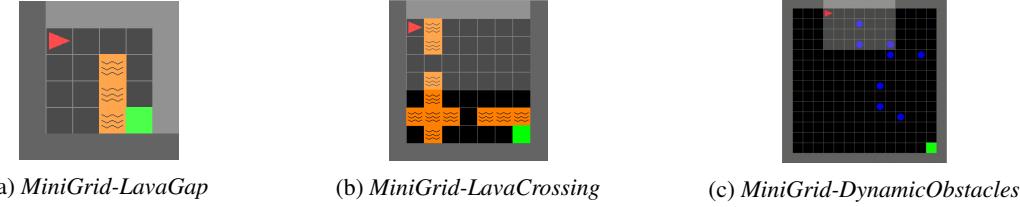


Figure 3: Screenshots of the environments used.

699 pixel representation of the observation, we use a shared representation for the value function and  
 700 policy, where the first component is a convolutional network, implemented as in Zhang [2018]. The  
 701 hyper-parameters of the neural representations are presented in Table 3.  
 702

Layer	Output	Func.
Conv1	16	ReLU
Conv2	32	ReLU
Conv3	64	ReLU

Table 3: Shared Observation Layers

703 The actor and critic layers, for both algorithms, are a fully connected layer with 64 features as input  
 704 and the corresponding output. We used in all cases an Adam optimiser. We optimised the parameters  
 705 for each (vanilla) algorithm through a quick parameter search, and apply the same parameters for the  
 706 Lexicographically Robust versions.

	LavaGap	LavaCrossing	DynamicObstacles
Parallel Envs	16	16	16
Steps	$2 \cdot 10^6$	$2 \cdot 10^6$	$8 \times 10^6$
$\gamma$	0.99	0.99	0.98
$\alpha$	0.00176	0.00176	0.00181
$\epsilon$ (Adam)	$10^{-8}$	$10^{-8}$	$10^{-8}$
Grad. Clip	0.9	0.9	0.5
Gae	0.95	0.95	0.95
Rollout	64	64	64
E. Coeff	0.01	0.014	0.011
V. Coeff	0.05	0.05	0.88

Table 4: A2C Parameters

	LavaGap	LavaCrossing	DynamicObstacles
Parallel Envs	8	8	8
Steps	$6 \cdot 10^6$	$2 \cdot 10^6$	$8 \times 10^5$
$\gamma$	0.95	0.99	0.97
$\alpha$	0.001	0.001	0.001
$\epsilon$ (Adam)	$10^{-8}$	$10^{-8}$	$10^{-8}$
Grad. Clip	1	1	0.1
Ratio Clip	0.2	0.2	0.2
Gae	0.95	0.95	0.95
Rollout	256	512	256
Epochs	10	10	10
E. Coeff	0	0.1	0.01

Table 5: PPO Parameters

707 For the implementation of the LRPG versions of the algorithms, in all cases we allow the algorithm to  
 708 iterate for 1/3 of the total steps before starting to compute the robustness objectives. In other words,  
 709 we use  $\hat{K}(\theta) = K_1(\theta)$  until  $t = \frac{1}{3} \text{max\_steps}$ , and from this point we resume the lexicographic  
 710 robustness computation as described in Algorithm 1. This is due to the structure of the environments  
 711 simulated. The rewards (and in particular the positive rewards) are very sparse in the environments  
 712 considered. Therefore, when computing the policy gradient steps, the loss for the primary objective  
 713 is practically zero until the environment is successfully solved at least once. If we implement the  
 714 combined lexicographic loss from the first time step, many times the algorithm would converge to a  
 715 (constant) policy without exploring for enough steps, leading to convergence towards a maximally  
 716 robust policy that does not solve the environment.

717 **Noise Kernels.** We consider two types of noise; a normal distributed noise  $\tilde{T}^g$  and a uniform  
 718 distributed noise  $\tilde{T}^u$ . For the environments LavaGap and DynamicObstacles, the kernel  $\tilde{T}^u$  produces  
 719 a disturbed state  $\tilde{x} = x + \xi$  where  $\|\xi\|_\infty \leq 2$ , and for LavaCrossing  $\|\xi\|_\infty \leq 1.5$ . The normal  
 720 distributed noise is in all cases  $\mathcal{N}(0, 0.5)$ . The maximum norm of the noise is quite large, but this  
 721 is due to the structure of the observations in these environments. The pixel values are encoded as  
 722 integers 0 – 9, where each integer represents a different feature in the environment (empty space,  
 723 doors, lava, obstacle, goal...). Therefore, any noise  $\|\xi\|_\infty \leq 0.5$  would most likely not be enough to  
 724 confuse the agent. On the other hand, too large noise signals are unrealistic and produce pathological  
 725 environments. All the policies are then tested against two “true” noise kernels,  $T_1 = \tilde{T}^u$  and  $T_2 = \tilde{T}^g$ .  
 726 The main reason for this is to test both the scenarios where we assume a *wrong* noise kernel, and the  
 727 case where we are training the agents with the correct kernel.

728 **Comparison with SA-PPO.** One of the baselines included is the State-Adversarial PPO algorithm  
 729 proposed in Zhang et al. [2020]. The implementation includes an extra parameter that multiplies the  
 730 regularisation objective,  $k_{ppo}$ . Since we were not able to find indications on the best parameter for  
 731 discrete action environments, we implemented  $k_{ppo} \in \{0.1, 1, 2\}$  and picked the best result for each  
 732 entry in Table 1. Larger values seemed to de-stabilise the learning in some cases. The rest of the  
 733 parameters are kept as in the vanilla PPO implementation.

### 734 C.1.1 Extended Results: Adversarial Disturbances

735 Even though we do not use an adversarial attacker or disturbance in our reasoning through this work,  
 736 we implemented a policy-based state-adversarial noise disturbance to test the benchmark algorithms  
 737 against, and evaluate how well each of the methods reacts to such adversarial disturbances.

738 **Adversarial Disturbance** We implement a bounded policy-based adversarial attack, where at each  
 739 state  $x$  we maximise for the KL divergence between the disturbed and undisturbed state, such that the  
 740 adversarial operator is:

$$T_{adv}^\varepsilon(y | x) = 1 \implies y \in \operatorname{argmax}_{\tilde{x}} D_{KL}(\pi(x), \pi(\tilde{x})) \\ s.t. \|x - \tilde{x}\|_2 \leq \varepsilon.$$

PPO on MiniGrid Environments					A2C on MiniGrid Environments				
Noise	Vanilla	LR <sub>PPO</sub> ( $K_T^u$ )	LR <sub>PPO</sub> ( $K_T^g$ )	SA-PPO	Vanilla	LR <sub>A2C</sub> ( $K_T^u$ )	LR <sub>A2C</sub> ( $K_T^g$ )	LR <sub>A2C</sub> ( $K_D$ )	
<i>LavaGap</i>									
$\emptyset$	<b>0.95±0.003</b>	<b>0.95±0.075</b>	<b>0.95±0.101</b>	0.94±0.068	<b>0.94±0.004</b>	<b>0.94±0.005</b>	<b>0.94±0.003</b>	<b>0.94±0.006</b>	
$T_1$	0.80±0.041	<b>0.95±0.078</b>	0.93±0.124	0.88±0.064	0.83±0.061	<b>0.93±0.019</b>	0.89±0.032	0.91±0.088	
$T_2$	0.92±0.015	<b>0.95±0.052</b>	<b>0.95±0.094</b>	0.93±0.050	0.89±0.029	<b>0.94±0.008</b>	0.93±0.011	0.93±0.021	
$T^{0.5}$	0.56±0.194	<b>0.93±0.101</b>	0.91±0.076	0.90±0.123	0.92±0.034	<b>0.94±0.003</b>	<b>0.94±0.007</b>	0.93±0.015	
$T_1^{adv}$	0.20±0.243	<b>0.90±0.124</b>	0.68±0.190	<b>0.90±0.135</b>	0.75±0.123	<b>0.94±0.006</b>	0.92±0.038	0.88±0.084	
$T_2^{adv}$	0.01±0.051	0.71±0.251	0.21±0.357	<b>0.87±0.116</b>	0.27±0.119	<b>0.79±0.069</b>	0.68±0.127	0.56±0.249	
<i>LavaCrossing</i>									
$\emptyset$	<b>0.95±0.023</b>	0.93±0.050	0.93±0.018	0.88±0.091	0.91±0.024	0.91±0.063	0.90±0.017	<b>0.92±0.034</b>	
$T_1$	0.50±0.110	<b>0.92±0.053</b>	0.89±0.029	0.64±0.109	0.66±0.071	<b>0.78±0.111</b>	0.72±0.073	0.76±0.098	
$T_2$	0.84±0.061	<b>0.92±0.050</b>	<b>0.92±0.021</b>	0.85±0.094	0.78±0.054	0.83±0.105	0.86±0.029	<b>0.87±0.063</b>	
$T^{0.5}$	0.29±0.098	<b>0.91±0.081</b>	<b>0.91±0.054</b>	0.87±0.045	0.56±0.039	0.51±0.089	0.43±0.041	<b>0.68±0.126</b>	
$T_1^{adv}$	0.03±0.022	0.83±0.122	0.86±0.132	<b>0.87±0.059</b>	0.27±0.158	0.25±0.118	0.17±0.067	<b>0.43±0.060</b>	
$T_2^{adv}$	0.0±0.004	0.50±0.171	0.38±0.020	<b>0.82±0.072</b>	0.06±0.056	0.04±0.030	0.01±0.008	<b>0.09±0.060</b>	
<i>DynamicObstacles</i>									
$\emptyset$	<b>0.91±0.002</b>	<b>0.91±0.008</b>	<b>0.91±0.007</b>	<b>0.91±0.131</b>	<b>0.91±0.011</b>	0.88±0.020	0.89±0.009	<b>0.91±0.013</b>	
$T_1$	0.23±0.201	<b>0.77±0.102</b>	0.61±0.119	0.45±0.188	0.27±0.104	0.43±0.108	0.45±0.162	<b>0.56±0.270</b>	
$T_2$	0.50±0.117	<b>0.75±0.075</b>	0.70±0.072	0.68±0.490	0.45±0.086	0.53±0.109	0.52±0.161	<b>0.67±0.203</b>	
$T^{0.5}$	0.74±0.230	0.89±0.118	0.85±0.061	<b>0.90±0.142</b>	0.46±0.214	0.55±0.197	0.51±0.371	<b>0.62±0.249</b>	
$T_1^{adv}$	0.26±0.269	0.79±0.157	0.68±0.144	<b>0.84±0.150</b>	0.19±0.284	<b>0.35±0.197</b>	0.23±0.370	0.10±0.379	
$T_2^{adv}$	-0.49±0.312	0.51±0.234	0.33±0.202	<b>0.55±0.170</b>	-0.54±0.209	-0.21±0.192	-0.53±0.261	<b>-0.51±0.260</b>	

Table 6: Extended Reward Results.

741 The optimisation problem is solved at every point by using a Stochastic Gradient Langevin Dynamics  
742 (SGLD) optimiser. The results are presented in Table 6.

743 This type of adversarial attack with SGLD optimiser was proposed in Zhang et al. [2020]. As one can  
744 see, the adversarial disturbance is quite successful at severely lowering the obtained rewards in all  
745 scenarios. Additionally, as expected SA-PPO was the most effective at minimizing the disturbance  
746 effect (as it is trained with adversarial disturbances), although LRPG produces reasonably robust  
747 policies against this type of disturbances as well. At last, A2C appears to be much more sensitive to  
748 adversarial disturbances than PPO, indicating that the policies produced by PPO are by default more  
749 robust than A2C.

## 750 C.2 Continuous Control

751 The continuous control environments simulated are MountainCar, LunarLander and BipedalWalker.  
752 The policies used are in all cases MLP policies with ReLU gates and a (64, 64) feature extractor plus  
753 a fully connected layer to output the values and actions unless stated otherwise. The hyperparameters  
754 can be found in tables 7 and 8. The implementation is based on Stable Baselines 3 [Raffin et al.,  
755 2021] tuned algorithms.

756 **Noise Kernels.** We consider again two types of noise; a normal distributed noise  $\tilde{T}^g$  and a uniform  
757 distributed noise  $\tilde{T}^u$ . In all cases, algorithms are implemented with a state observation normalizer.  
758 That is, asymptotically all states will be observed to be in the set  $(-1, 1)$ . For this reason, the uniform  
759 noise is bounded at lower values than for the discrete control environments. For BipedalWalker  
760  $\|\xi\|_\infty \leq 0.05$  and for Lunarlander and MountainCar  $\|\xi\|_\infty \leq 0.1$ . Larger values were shown to  
761 destabilize learning.

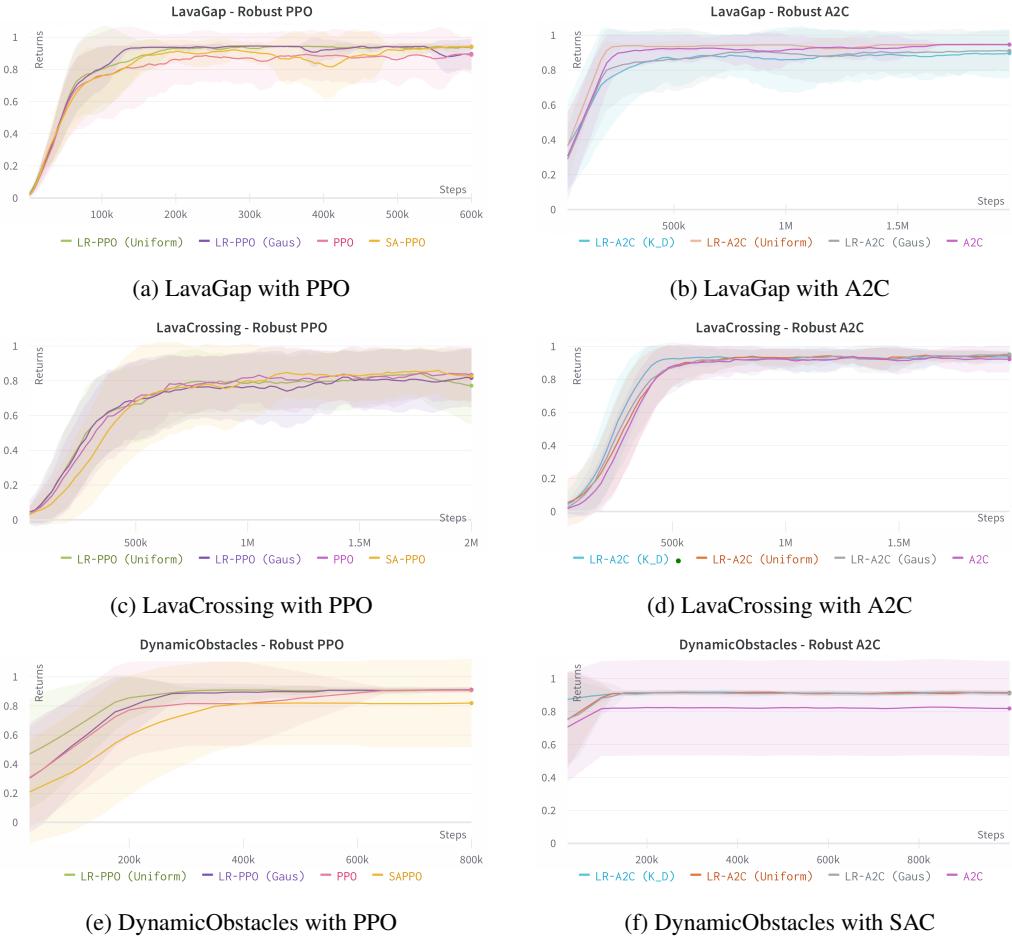


Figure 4: Learning Plots for Discrete Control Environments.

	MountainCarContinuous	LunarLanderContinuous	BipedalWalker-v3
Parallel Envs	1	16	32
Steps	$2 \times 10^4$	$1 \times 10^6$	$5 \times 10^6$
$\gamma$	0.9999	0.999	0.999
$\alpha$	$3 \times 10^{-4}$	$3 \times 10^{-4}$	$3 \times 10^{-4}$
Grad. Clip	5	0.5	0.5
Ratio Clip	0.2	0.2	0.18
Gae	0.9	0.98	0.95
Epochs	10	4	10
E. Coeff	0.00429	0.01	0

Table 7: PPO Parameters for Continuous Control

	MountainCarContinuous	LunarLanderContinuous	BipedalWalker-v3
Steps	$5 \times 10^4$	$5 \times 10^5$	$5 \times 10^5$
$\gamma$	0.9999	0.99	0.98
$\alpha$	$3 \times 10^{-4}$	$7.3 \times 10^{-4}$	$7.3 \times 10^{-4}$
$\tau$	0.01	0.01	0.01
Train Freq.	32	1	64
Grad. Steps	32	1	64
MLP Arch	(64,64)	(400,300)	(400,300)

Table 8: SAC Parameters for Continuous Control

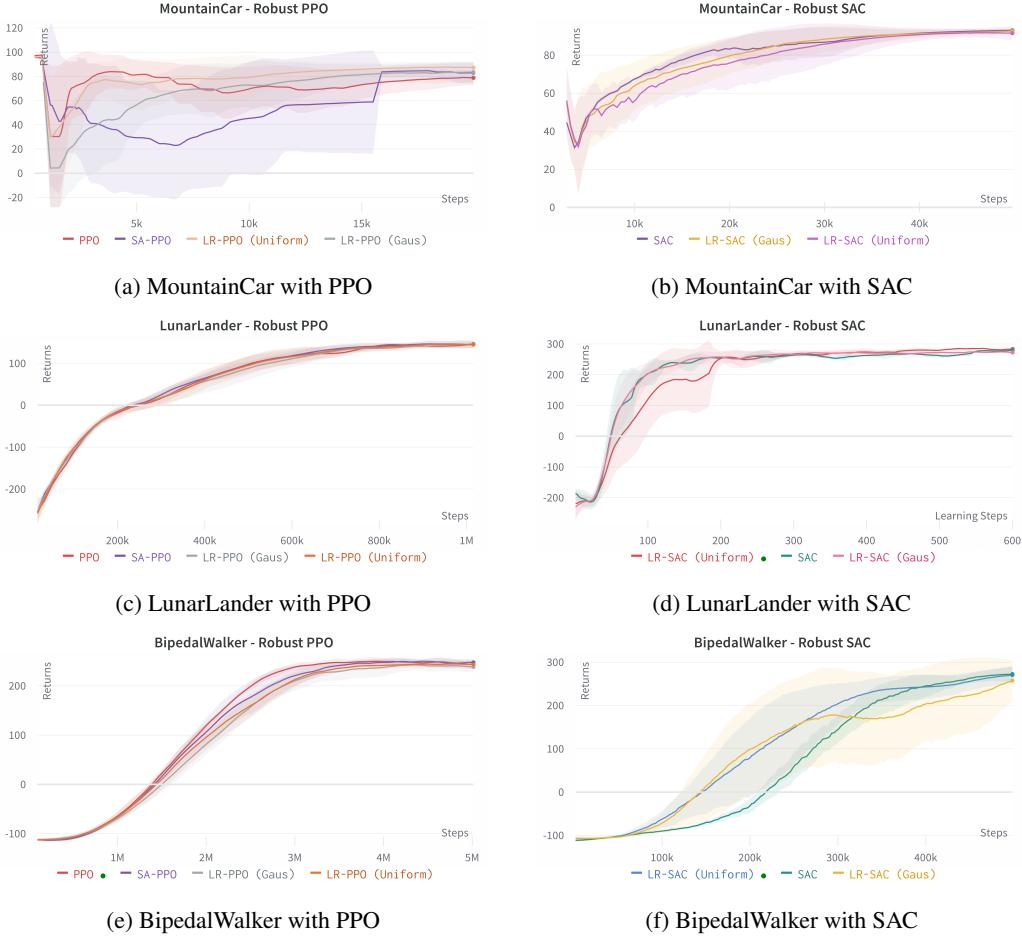


Figure 5: Learning Plots for Continuous Control Environments.

762 **Learning processes.** In general, learning was not severely affected by the LRPG scheme. However,  
763 it was shown to induce a larger variance in the trajectories observed, as seen in LunarLander with  
764 LR-SAC and BipedalWalker with LR-SAC.